



中国科学技术大学  
University of Science and Technology of China

# Software Architecture

SSE USTC   Qing Ding  
dingqing@ustc.edu.cn  
<http://staff.ustc.edu.cn/~dingqing>



# Serverless Architecture

## Function as a Service

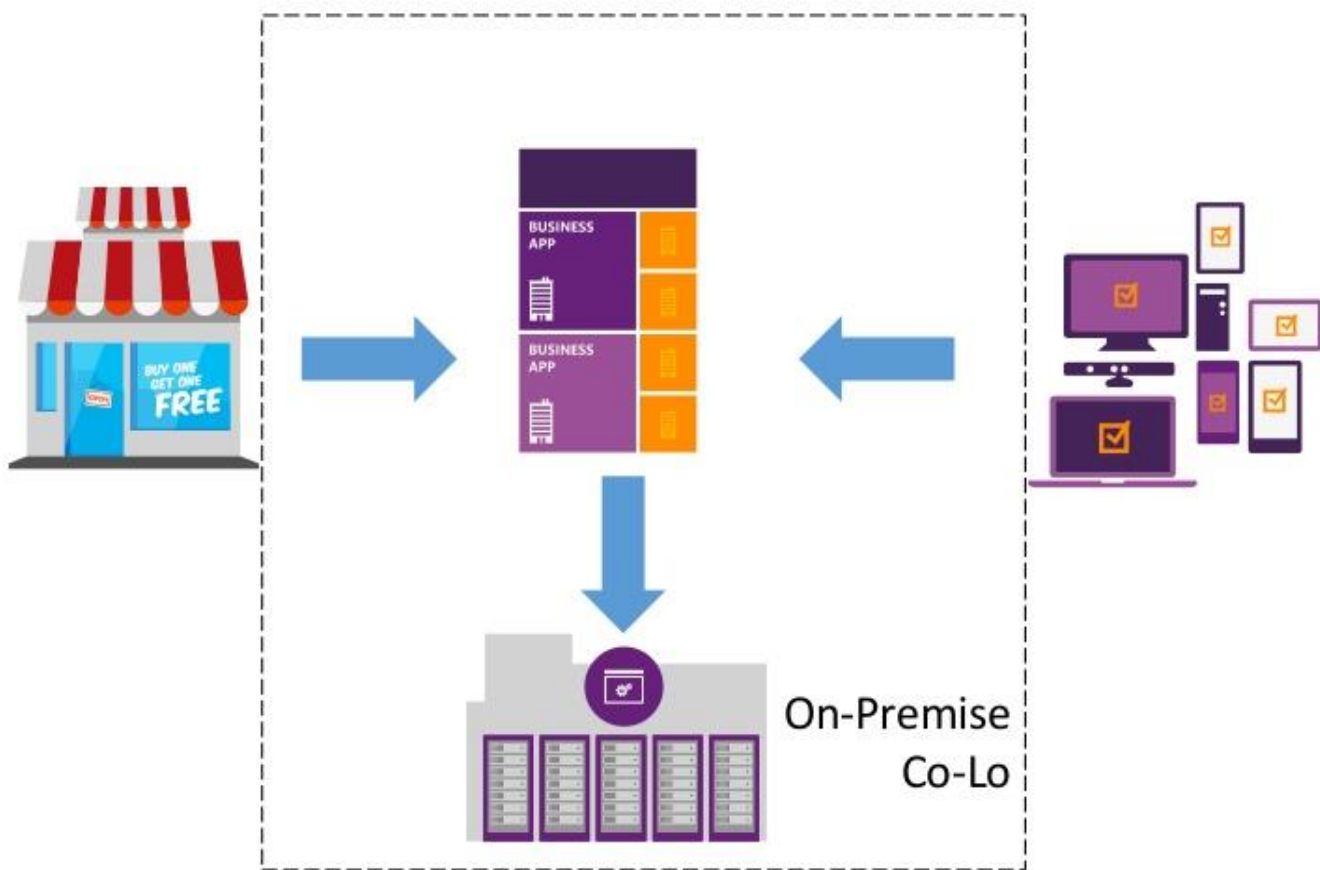
了解特点，不用知道具体细节



- Evolution of Serverless Computing
- *What is Serverless*
- Why is Serverless attractive
- AWS Architecture
- Apache OpenWhisk
- Azure Functions
- Future of Serverless: Research Challenges and Questions

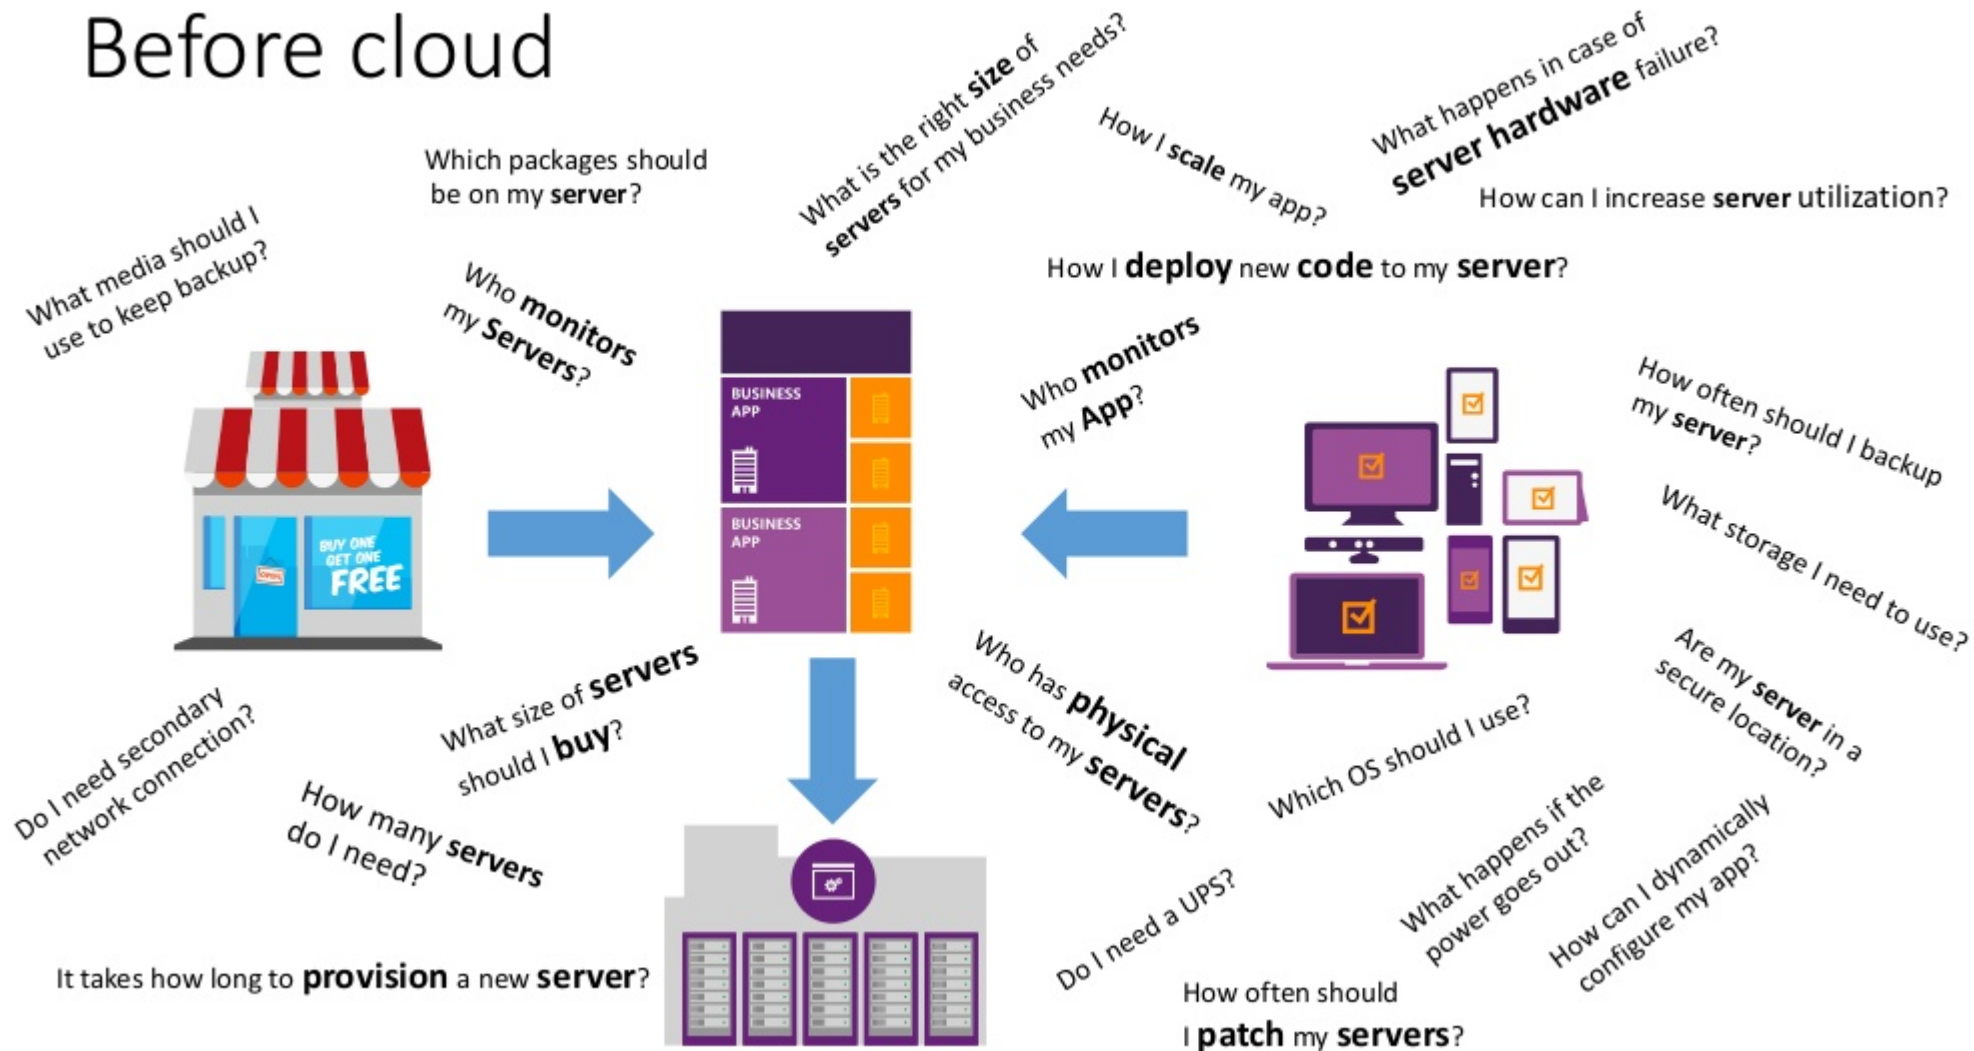


## Before cloud



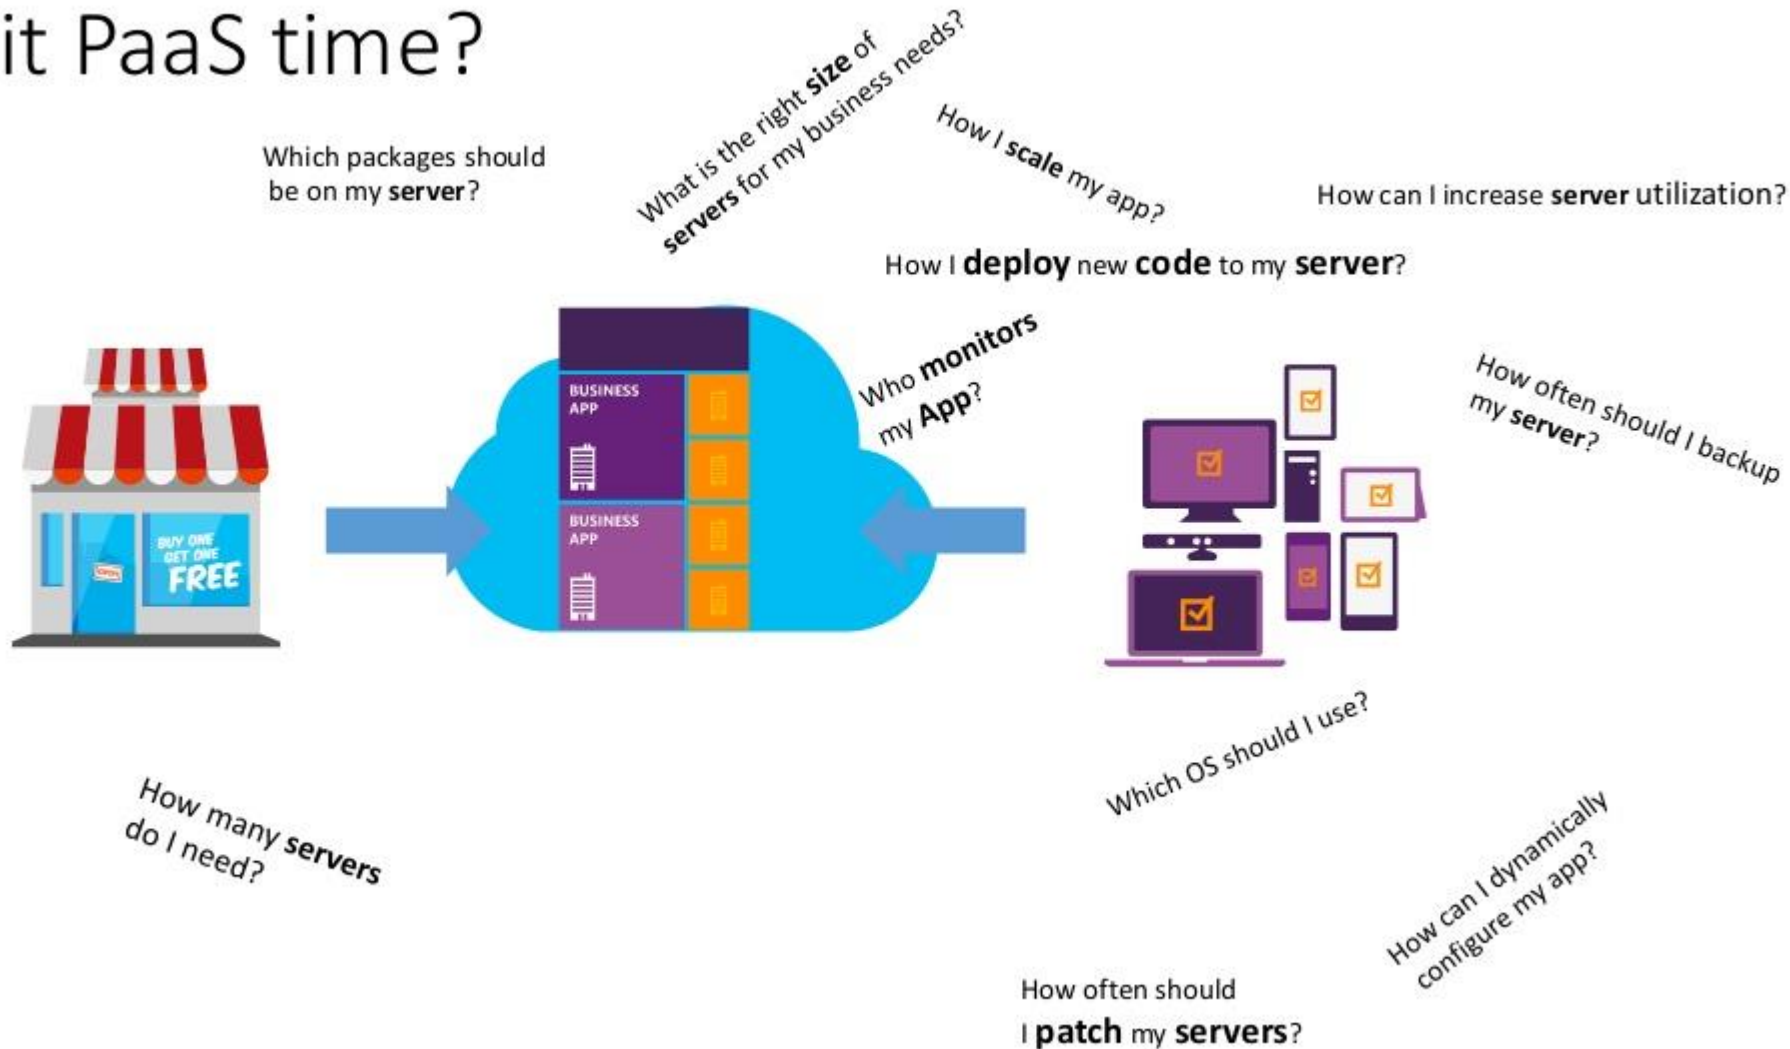


# Before cloud






# Is it PaaS time?







# Serverless . . .





# What is Serverless?



Abstraction  
of servers

服务抽象化



Event-driven/  
instant scale

事件驱动

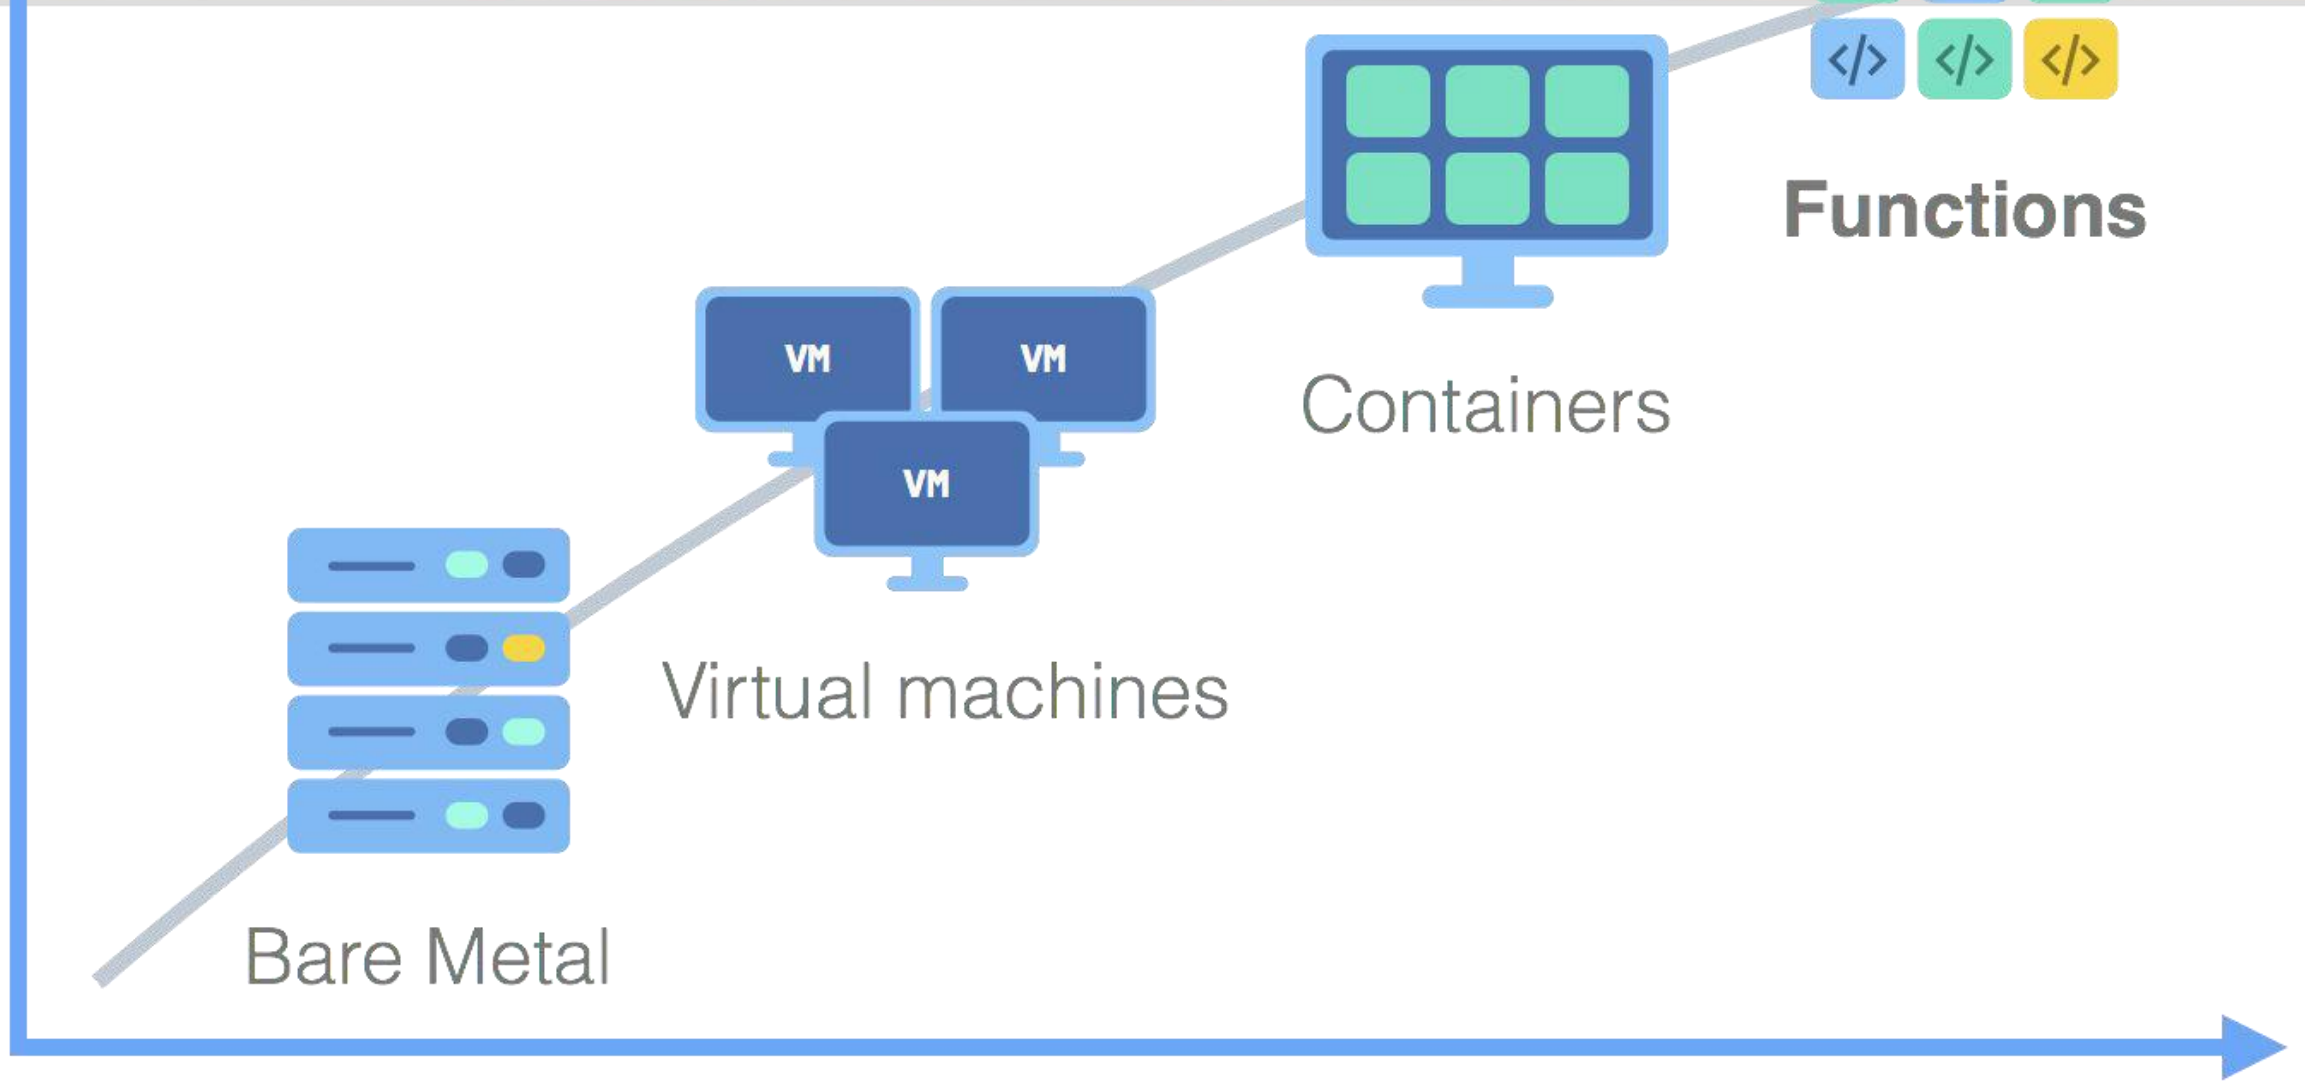


Sub-second  
billing





Increasing focus on business logic



**Functions**

Containers

Virtual machines

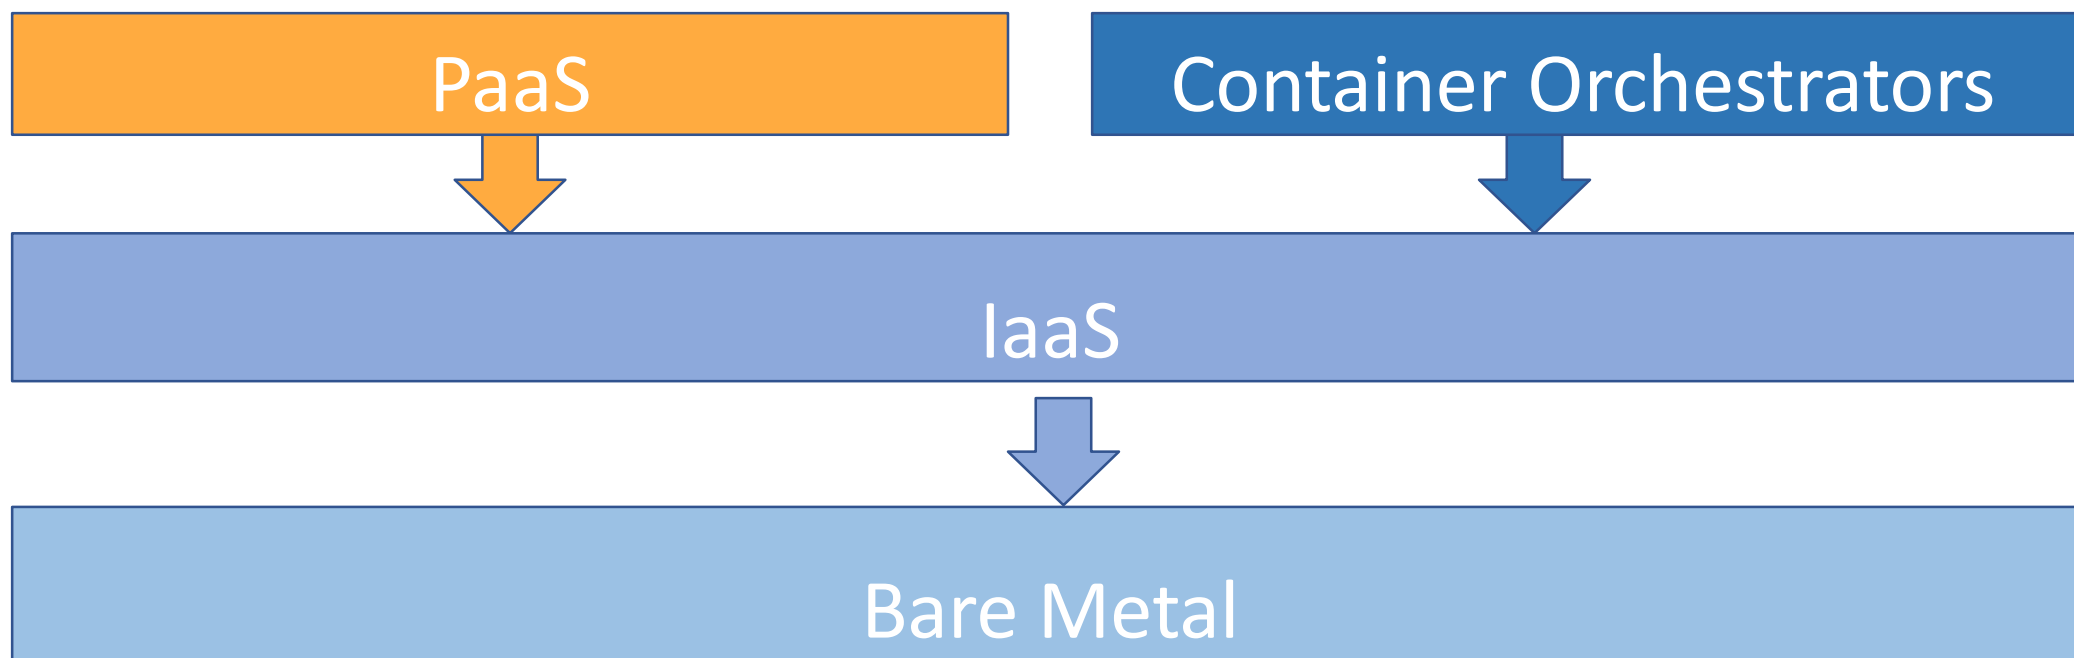
Bare Metal

Decreasing concern (and control) over stack implementation

# Evolution Of Serverless



中国科学技术大学  
University of Science and Technology of China





## Monolithic Application

Break-down into  
microservices




Explosion in number of  
containers / processes:

Make each micro service  
HA



Increase of infrastructure  
cost footprint

Protect against regional  
outages

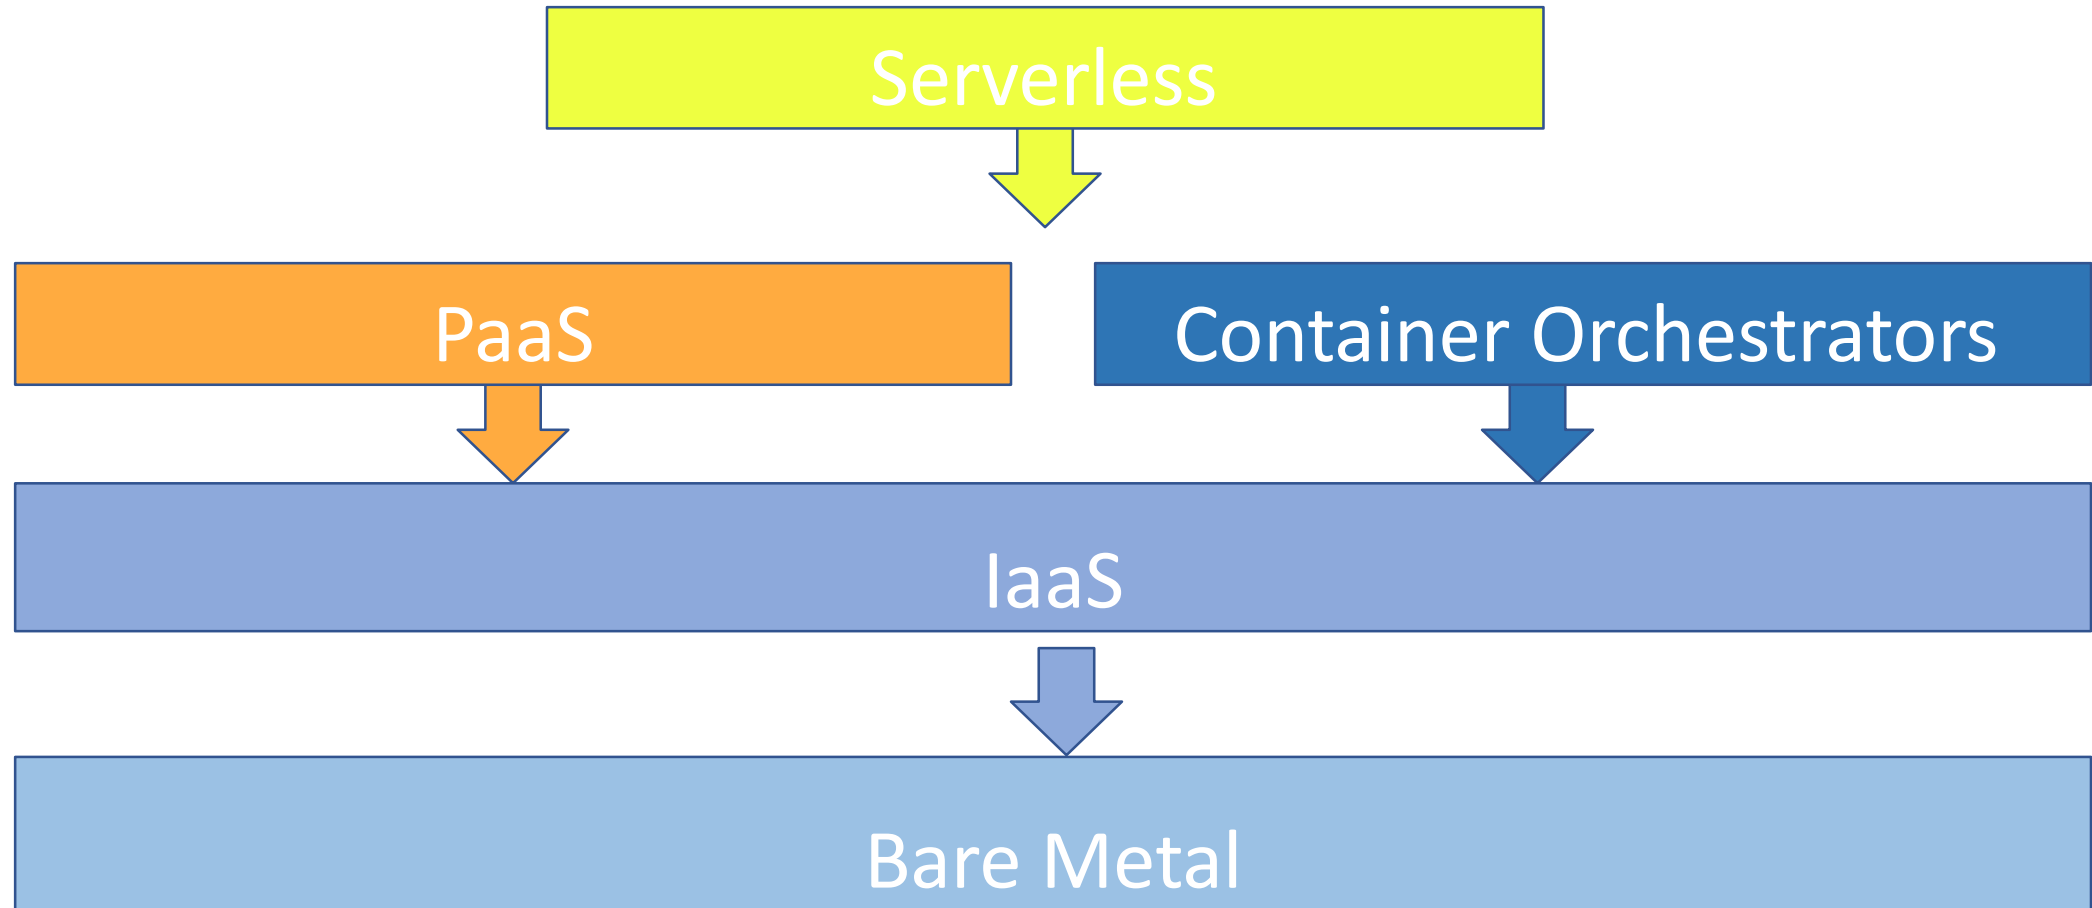


Increase of operational  
management cost and  
complexity

# Enter Serverless



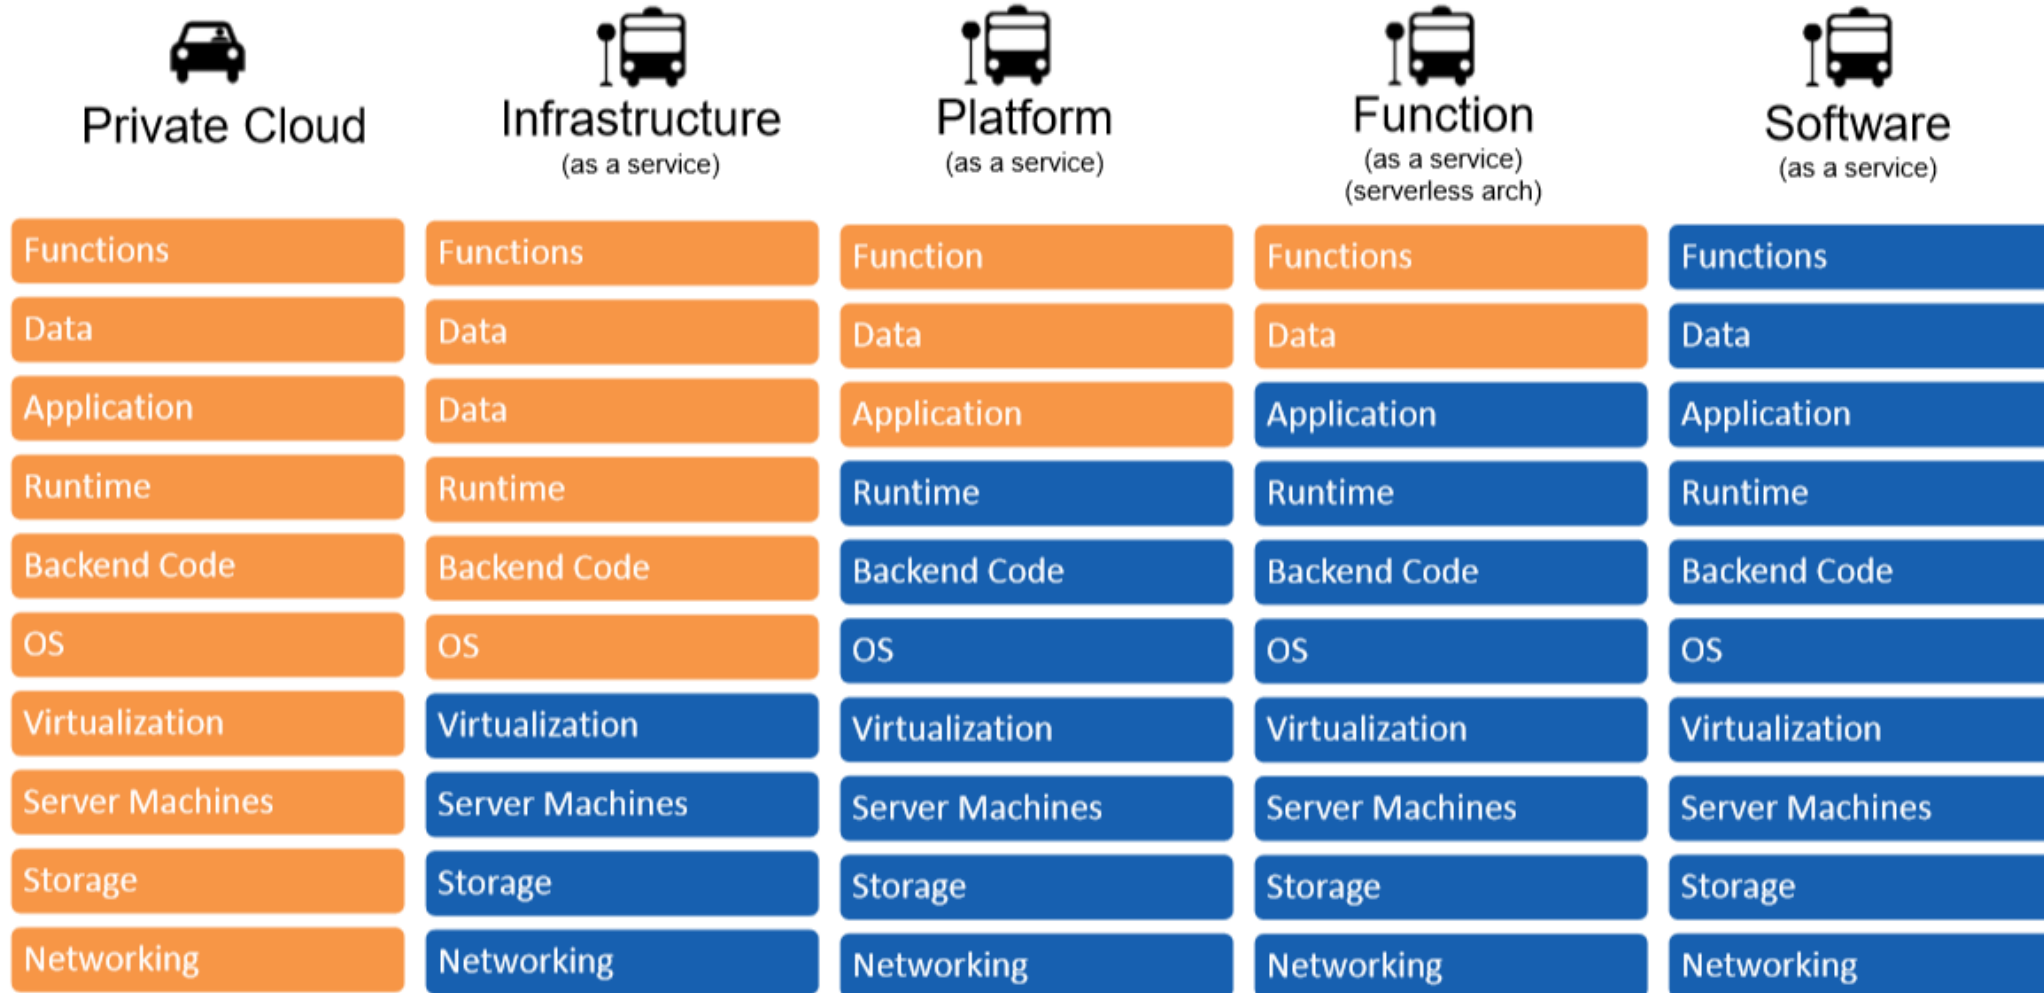
中国科学技术大学  
University of Science and Technology of China





# IaaS – PaaS, FaaS, SaaS



中国科学技术大学  
University of Science and Technology of China



 Public Cloud Provider - responsibility

 Application Writer - responsibility

[Awesome Vizualisation picked from : Ref : http://www.slideshare.net/manuel\\_silveyra/austin-cf-meetup-20150224/3](http://www.slideshare.net/manuel_silveyra/austin-cf-meetup-20150224/3)

PS: We expect Container as a Service term in 2017-18 too, there is a separate section on it later

# What is Serverless?



中国科学技术大学  
University of Science and Technology of China

a cloud-native platform

*for*

short-running, stateless computation

*and*

event-driven applications

*which*

scales up and down instantly and automatically

*and*

charges for actual usage at a millisecond granularity



# Server-less means no servers? Or worry-less about servers?



中国科学技术大学  
University of Science and Technology of China

Runs code **only** on-demand on  
a per-request basis

Serverless  
deployment &  
operations model



No servers




Just code

# What triggers code execution?




中国科学技术大学  
University of Science and Technology of China

Runs code **in response** to events




Event-programming  
model



# FaaS market is growing quickly




中国科学技术大学  
University of Science and Technology of China



# FaaS market is growing quickly



中国科学技术大学  
University of Science and Technology of China




# Google Search Trend over time



中国科学技术大学  
University of Science and Technology of China

Interest over time ?



# Why is Serverless attractive?



中国科学技术大学  
University of Science and Technology of China

- Making app development & ops dramatically faster, cheaper, easier
- Drives infrastructure cost savings

	On-prem	VMs	Containers	Serverless
Time to provision	Weeks-months	Minutes	Seconds-Minutes	Milliseconds
Utilization	Low	High	Higher	Highest
Charging granularity	CapEx	Hours	Minutes	Blocks of milliseconds



# Key factors for infrastructure cost savings



中国科学技术大学  
University of Science and Technology of China

	Traditional models (CF, containers, VMs)	Serverless
High Availability	At least 2-3 instances of everything	No incremental infrastructure
Multi-region deployment	One deployment per region	No incremental infrastructure
Cover delta between short (<10s) load spikes and valleys (vs average)	~2x of average load	No incremental infrastructure
Example incremental costs	2 instances x 2 regions x 2 = 8x	1x


# Benefits of Serverless Computing



中国科学技术大学  
University of Science and Technology of China



- No Servers to Manage
- Continuous Scaling
- Dynamic allocation of resources
- Avoid overallocation of resources
- Never Pay for Idle: pay-per-usage






## Benefits of Serverless?



Reduced  
DevOps



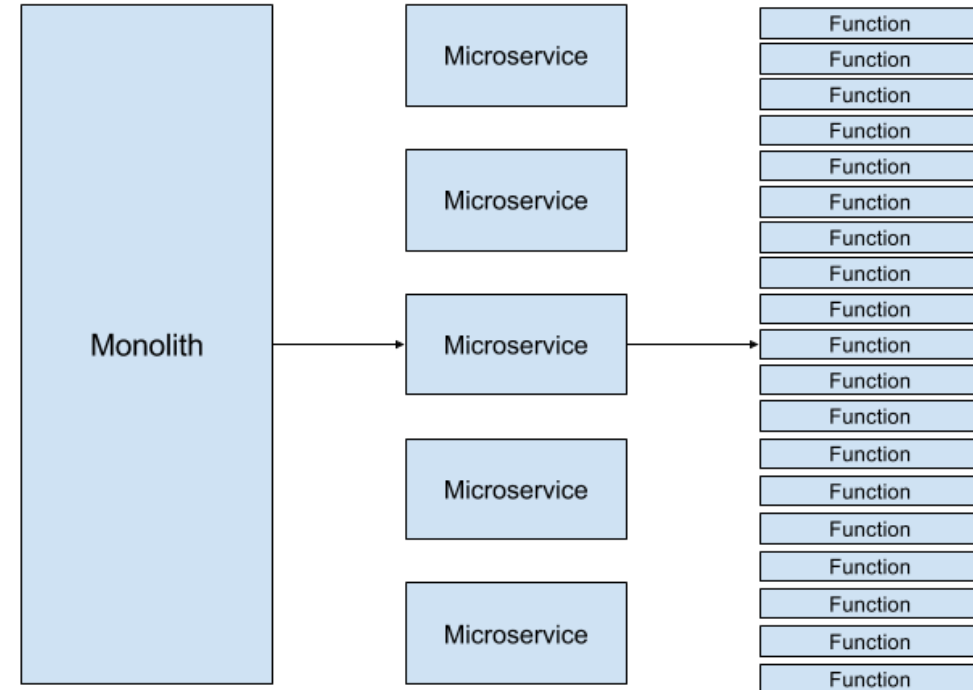
Focus on  
Business  
Logic



Reduced Time  
To Market

## Microservices

- Smaller-grained services
- Specified Functions
- Defined Capabilities



- Event handler
- Serverless back ends
- Data processing



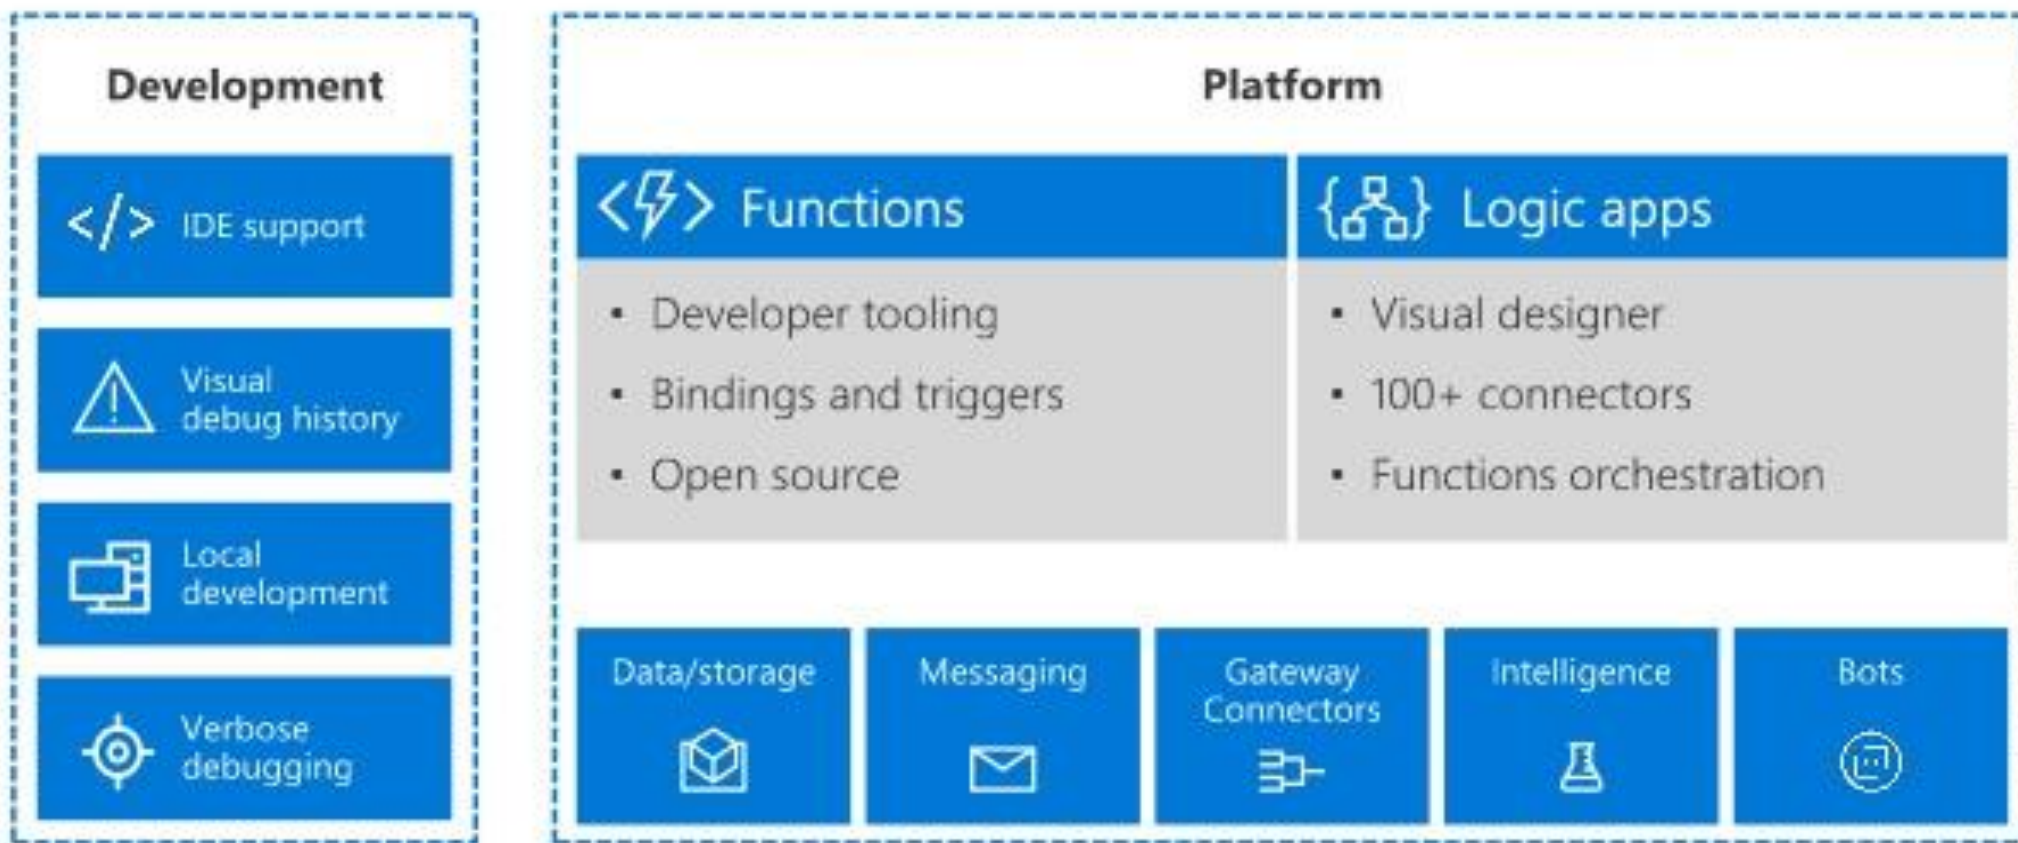
**“Serverless function”** or more accurately “Functions as a Service”

- **Principles of FaaS:**
- Complete abstraction of servers away from the developer
- Billing based on consumption and executions, not server instance sizes
- Services that are event-driven and instantaneously scalable






# Serverless application platform components



# What is Serverless good for?



中国科学技术大学  
University of Science and Technology of China

Serverless is **good** for  
*short-running stateless  
event-driven*

短时、无状态、事件驱动的



Microservices Mobile



Backends



Bots, ML Inferencing IoT



Modest Stream Processing



Service integration



Serverless is **not good**  
for

*long-running stateful  
number crunching*

长时、有状态、数值计算的



Databases



Deep Learning Training



Heavy-Duty Stream Analytics



Spark/Hadoop Analytics Numerical



Simulation




Video Streaming



# Current Platforms for Serverless



中国科学技术大学  
University of Science and Technology of China



AWS Lambda

OpenLambda



Azure Functions



Kubernetes



Red-Hat

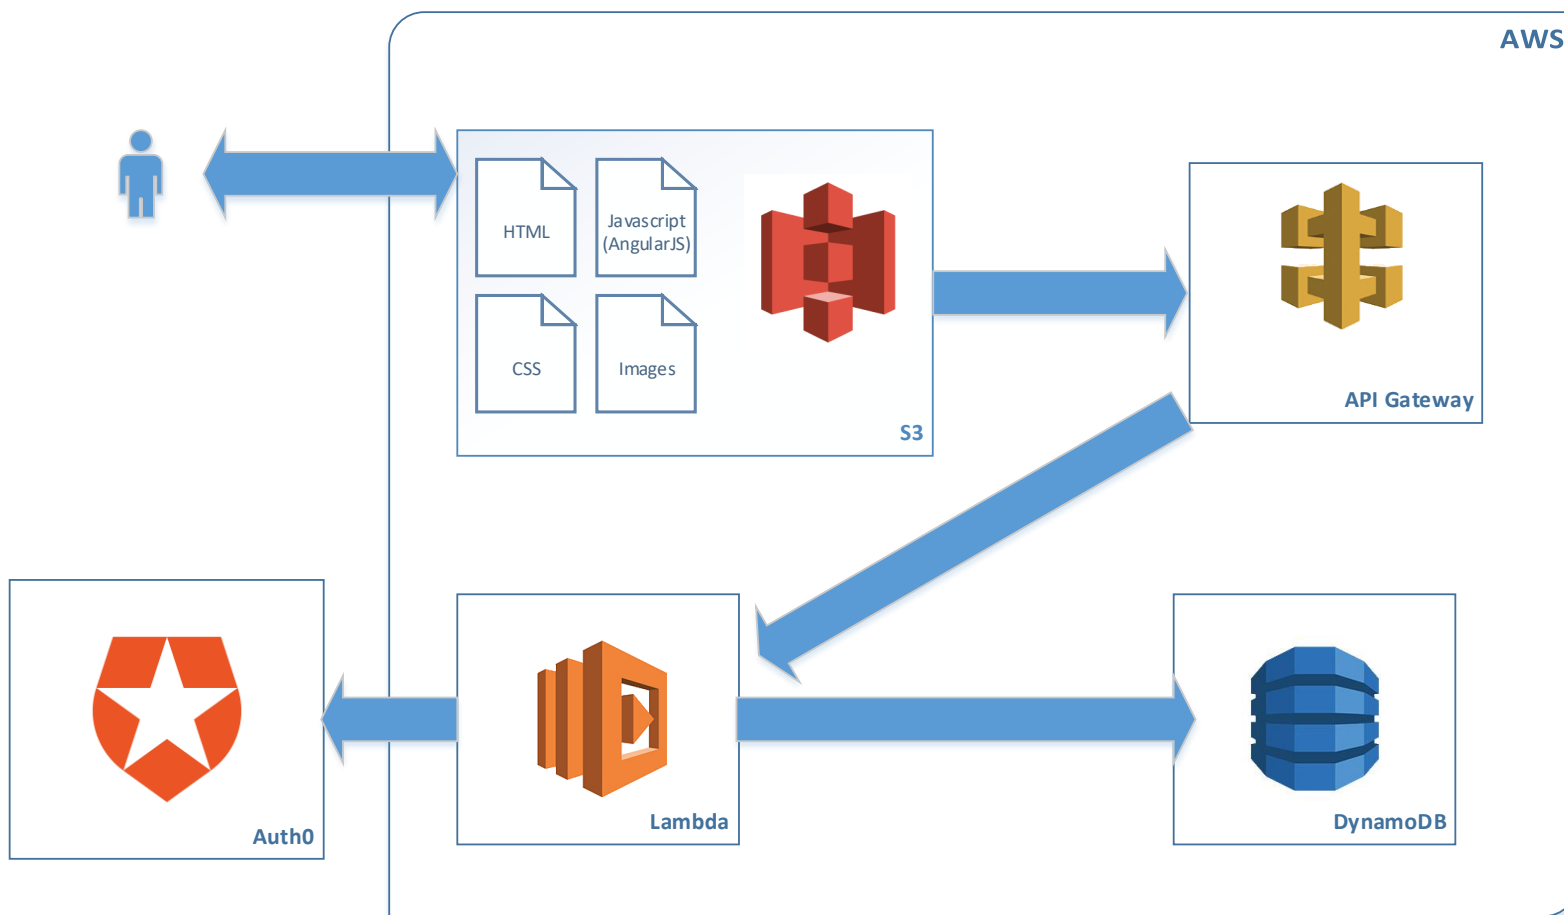


Google Functions

# AWS Architecture



中国科学技术大学  
University of Science and Technology of China





- Cloud document storage
  - Hosts static web resources
  - Can be configured to host public websites
    - HTML, CSS, JavaScript, images
    - Supports CNAME aliases
    - Supports .htaccess style configuration for URL rewrite and redirect
    - Supports SSL for custom domain names
  - Billed for storage and for transfer



- Acts as the ‘front door’ to the application
- Handles authorization and access control
- Exposes Lambda functions to your front-end application code
- Billed for API calls



- Run code without servers
  - Supports Node.js, Python, Java, C#, and Go.
  - Can use existing libraries
- Upload functions, configure triggers
- Works closely with API gateway via proxy integration
- Business logic and data access code goes here
- Billed for compute cycles (function run time)



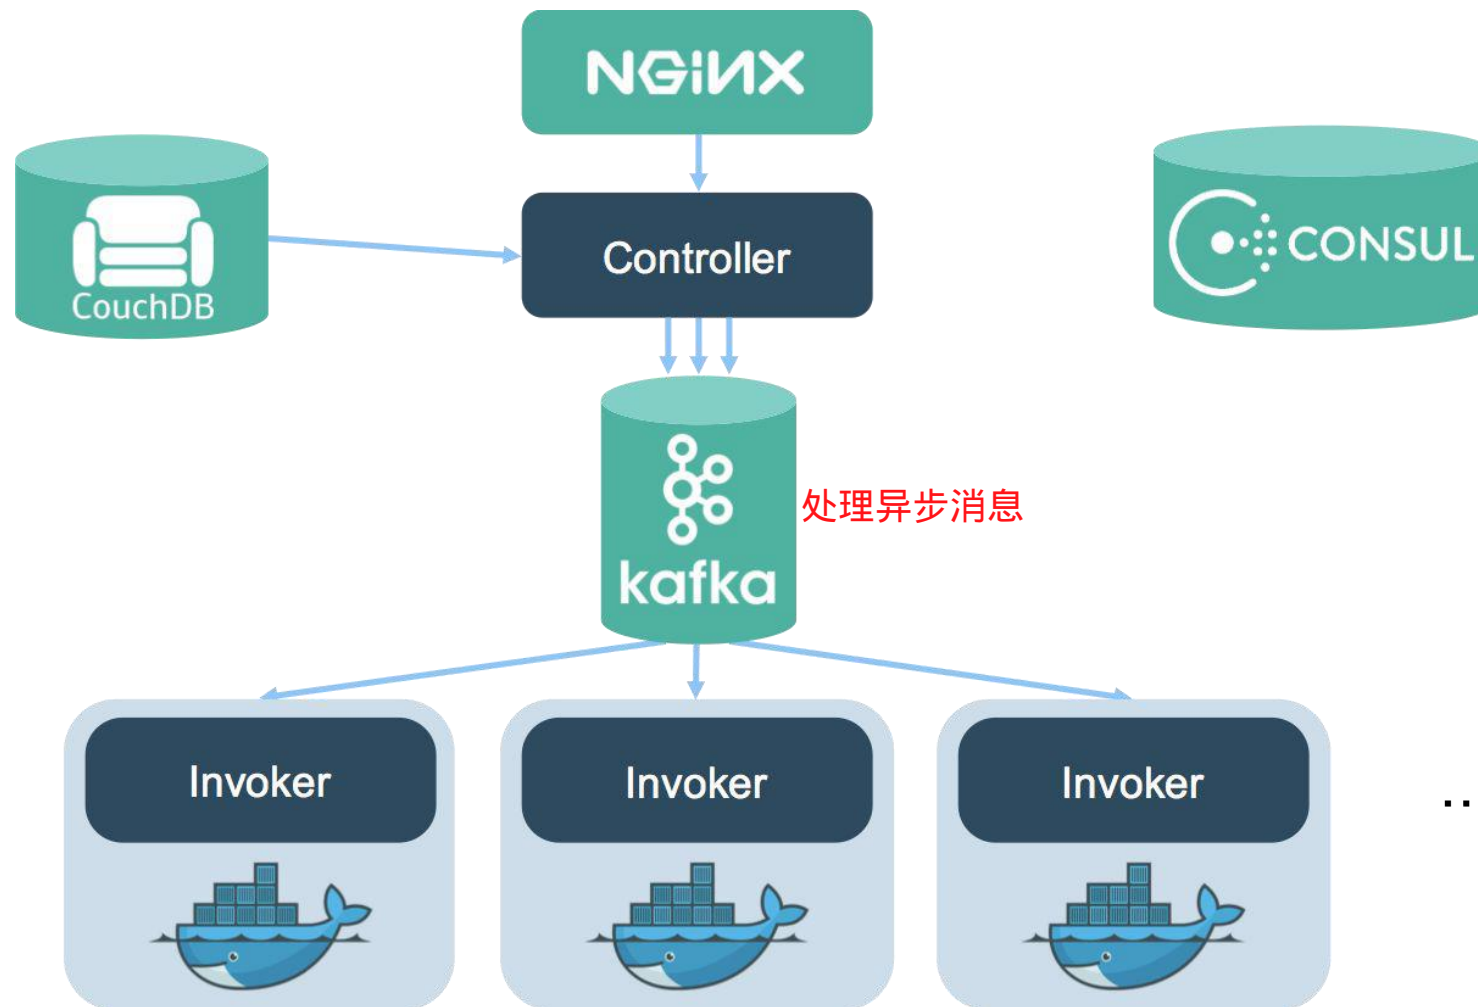
- Document database in the cloud
  - Non-relational (NoSQL)
  - Fast and scalable
- Built in security, backup, restore, caching
- Billed for storage and for transfer



# Apache OpenWhisk Serverless Architecture



科学技术大学  
University of Science and Technology of China




# Apache OpenWhisk:

## High-level serverless programming model

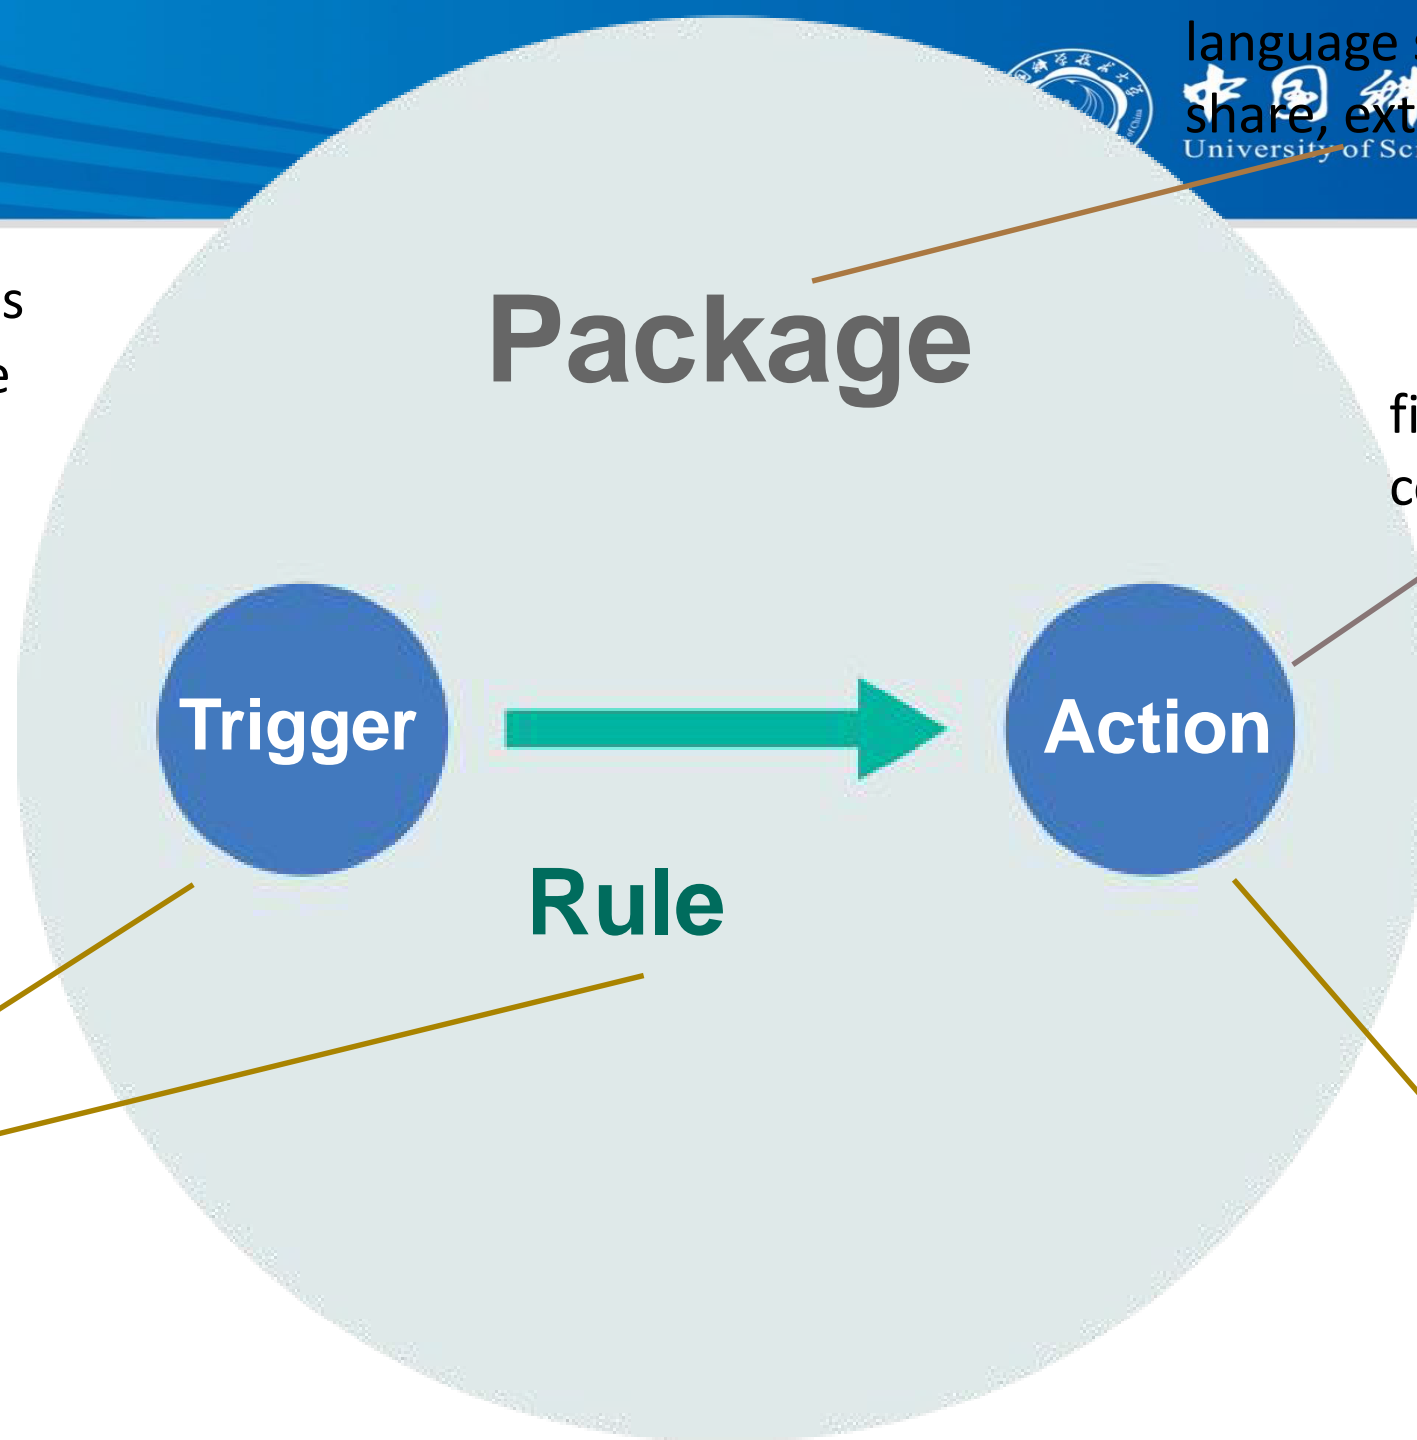


中国科学技术大学  
University of Science and Technology of China





all constructs first-class  
— powerful extensible  
language



first-class functions  
compose via sequences

first-class  
event-driven  
programming  
constructs

docker  
containers as  
actions



A

**Action:** a stateless function  
(event handler)



A

## Action

```
function main(params) { console.log( "Hello  
    " + params.name);  
return { msg: "Goodbye    " + params.name) };  
}
```




A

## Action: Python

```
def lambda_handler(event, context):  
    print("hello world")
```






A

## Action

```
func main(params:[String:Any]) -> [String:Any]
{
    var reply = [String:Any] ()
    if let name = params[ "name" ] as? String
    {
        print( "Hello \(name)" )
        reply[ "msg" ] = "Goodbye \(name)"
    }
    return reply
}
```



## A Action: sequence





# Trigger: a class of events (feed)



# AWS Lambda Trigger Sources



中国科学技术大学  
University of Science and Technology of China

## DATA STORES



Amazon S3



Amazon  
DynamoDB



Amazon  
Kinesis



Amazon  
Cognito

## ENDPOINTS



Amazon  
Alexa



Amazon  
API Gateway



AWS IoT

## CONFIGURATION REPOSITORIES



AWS  
CloudFormation



AWS  
CloudTrail



AWS  
CodeCommit



Amazon  
CloudWatch

## EVENT/MESSAGE SERVICES



Amazon  
SES



Amazon SNS



Cron events



R

Rule

Trigger

Action




# Apache OpenWhisk: Step 1. Entering the system



中国科学技术大学  
University of Science and Technology of China

POST /api/v1/namespaces/myNamespace/actions/myAction







Master VM

**controller**

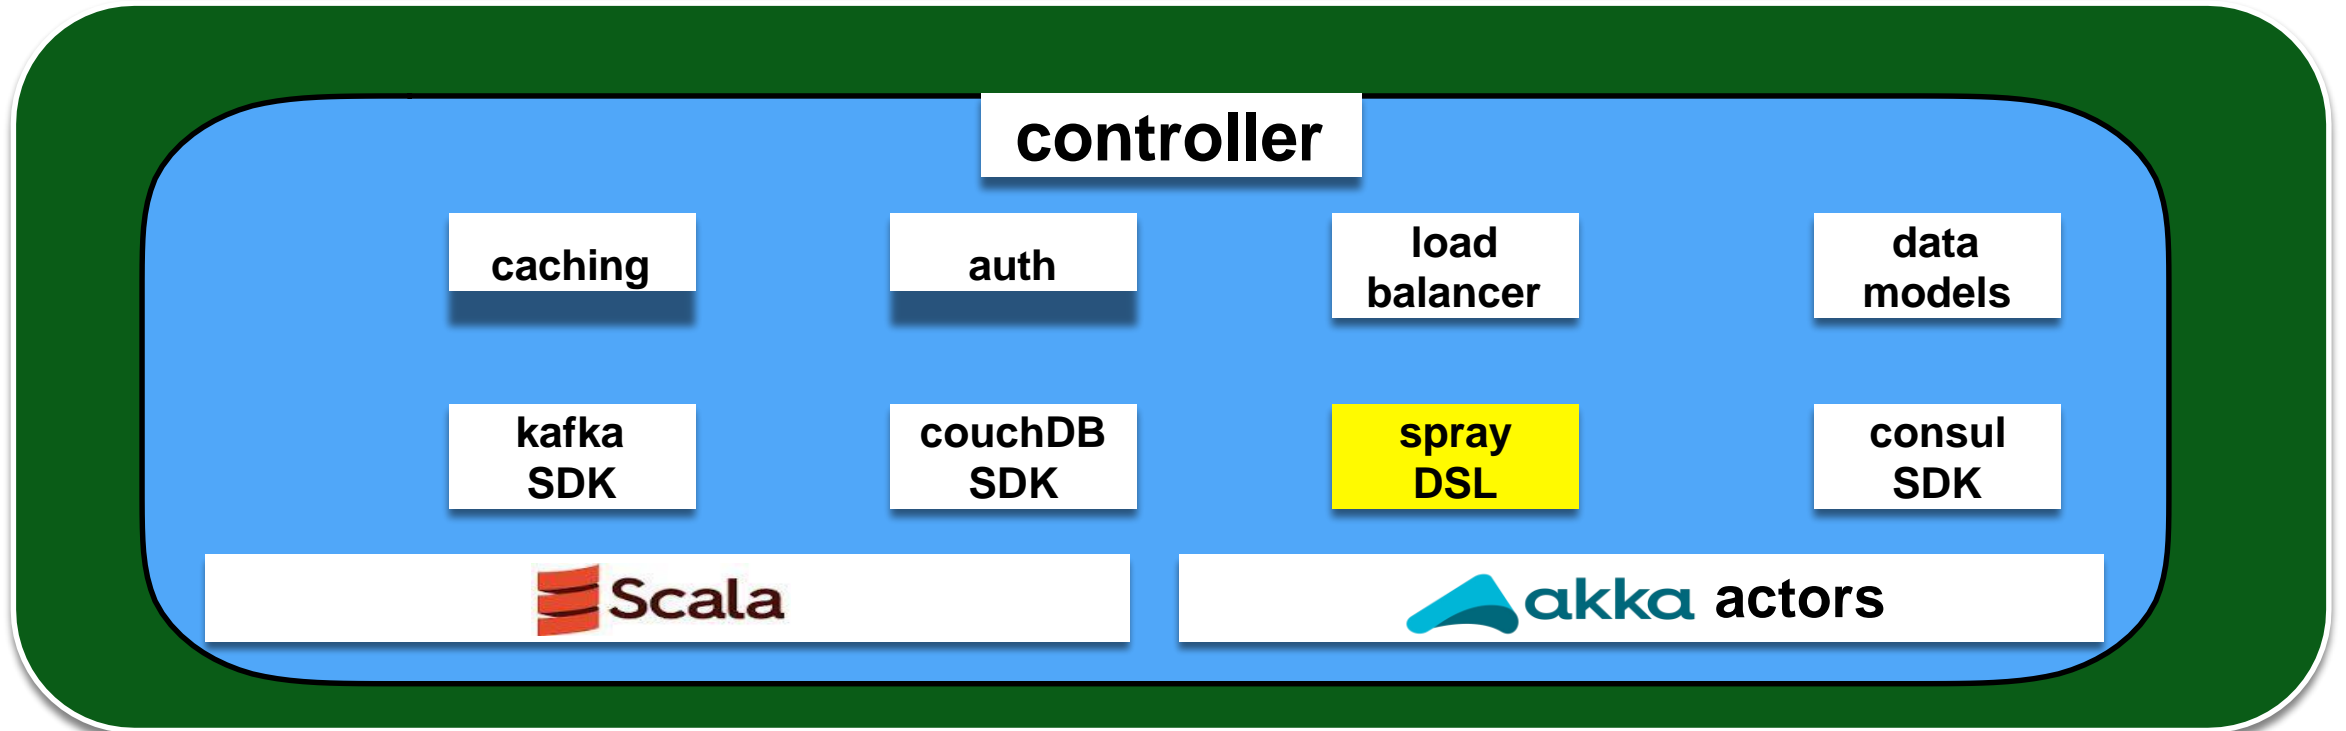


# Apache OpenWhisk: Step 2. Handle the request



中国科学技术大学  
University of Science and Technology of China

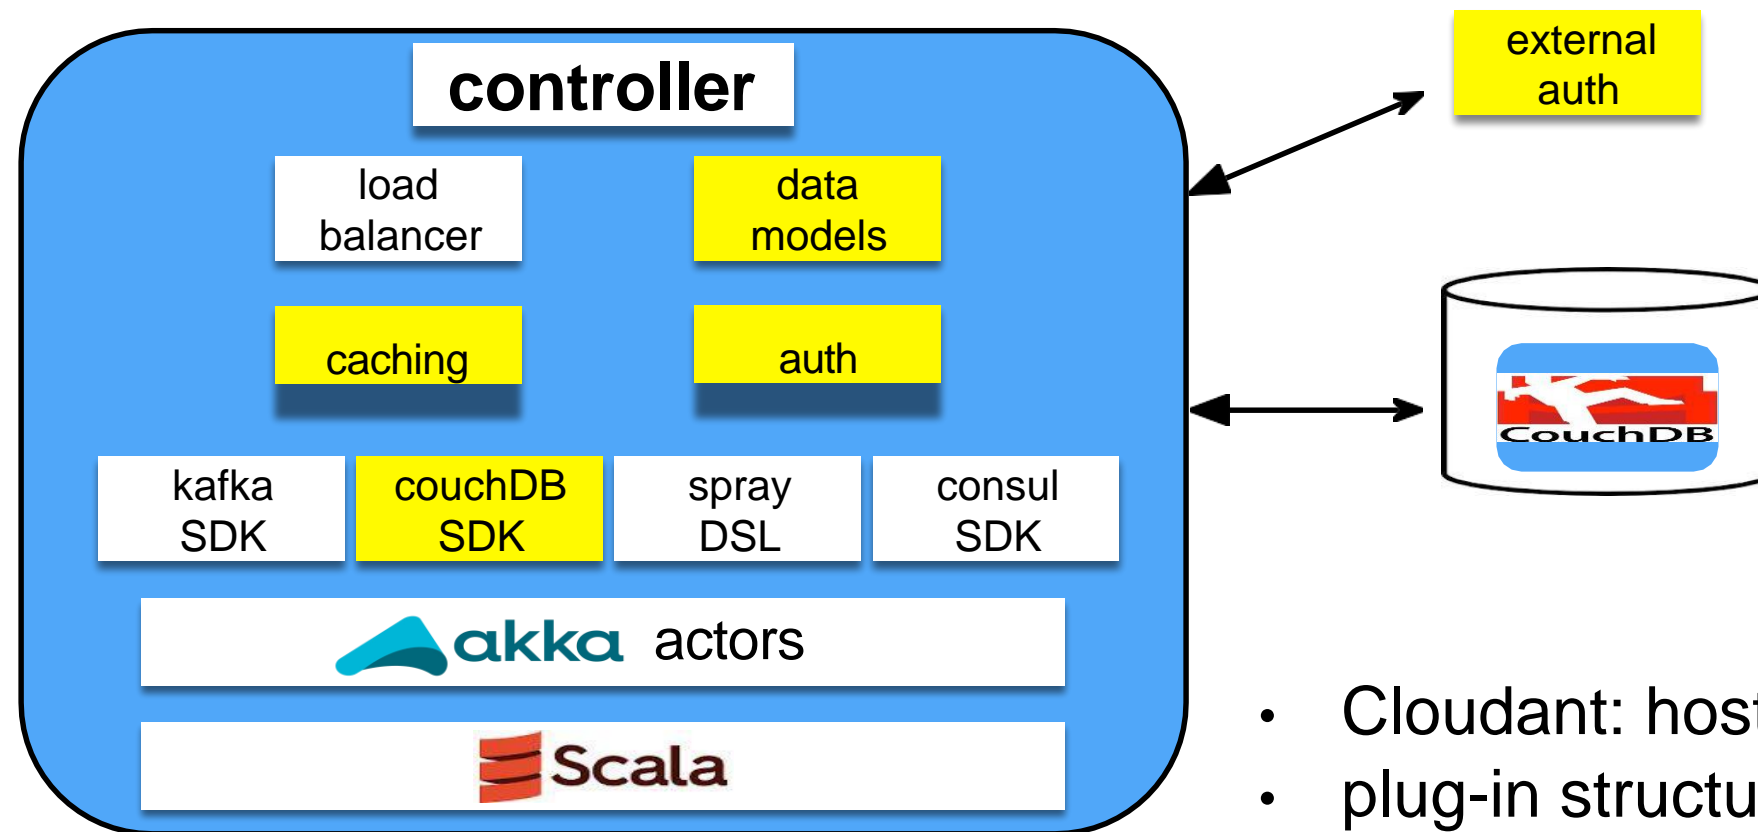
Master VM



# Apache OpenWhisk: Step 3. Authentication Authorization



中国科学技术大学  
University of Science and Technology of China



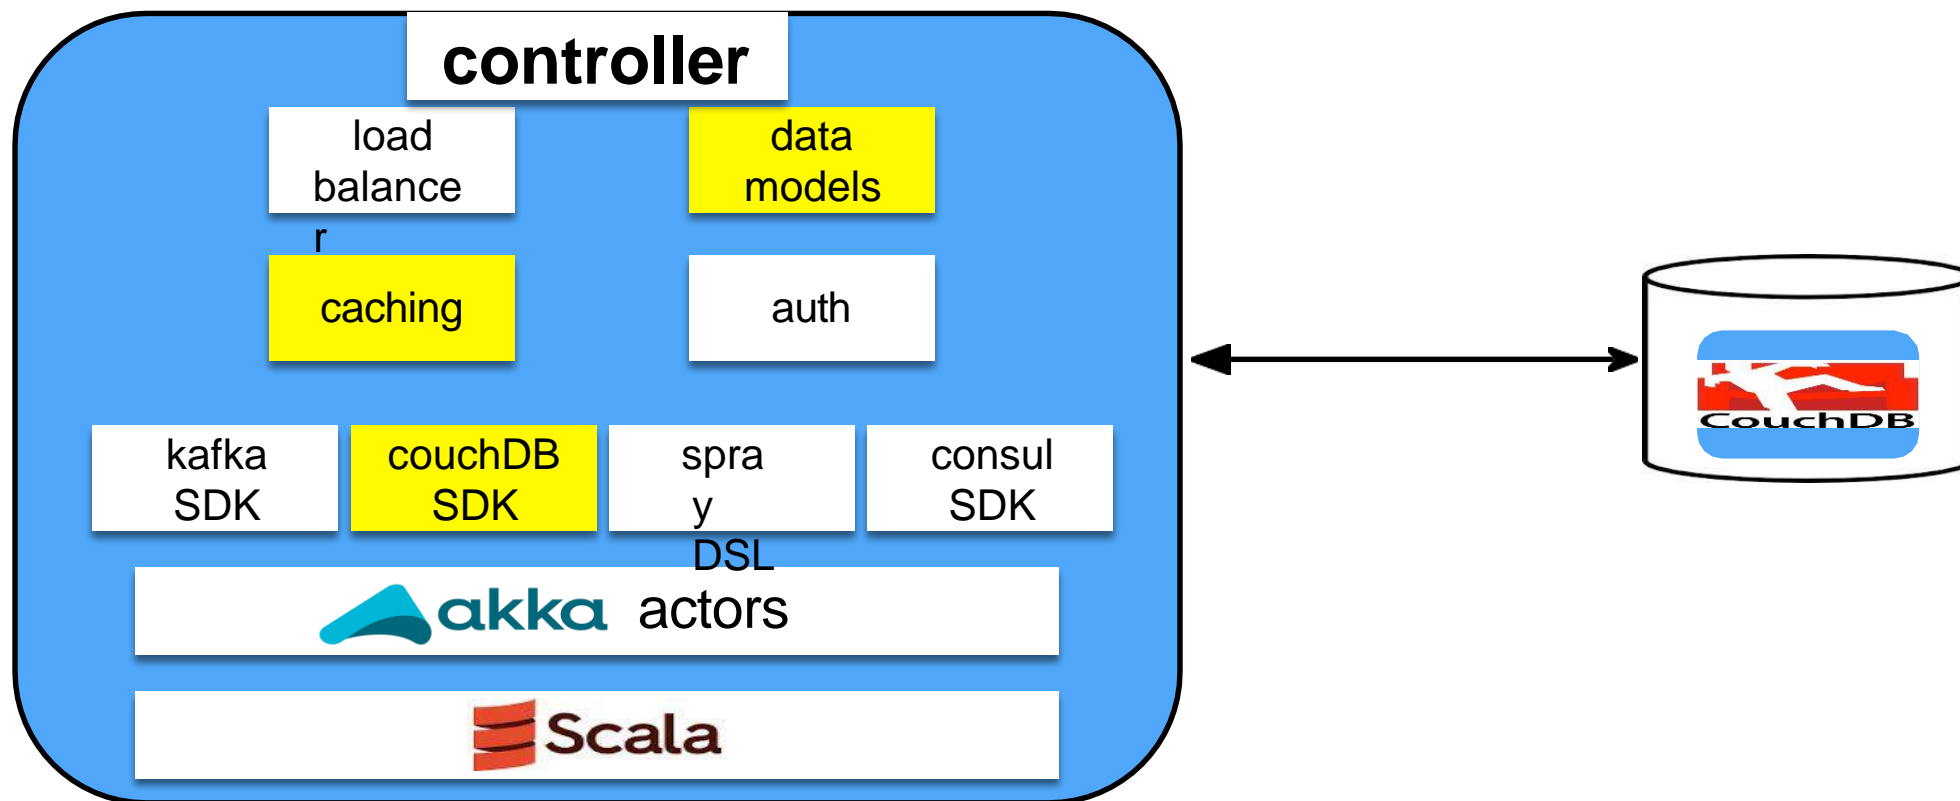
- Cloudfant: hosted CouchDB
- plug-in structure for custom authentication module

# Apache OpenWhisk: Step 4. Get the action


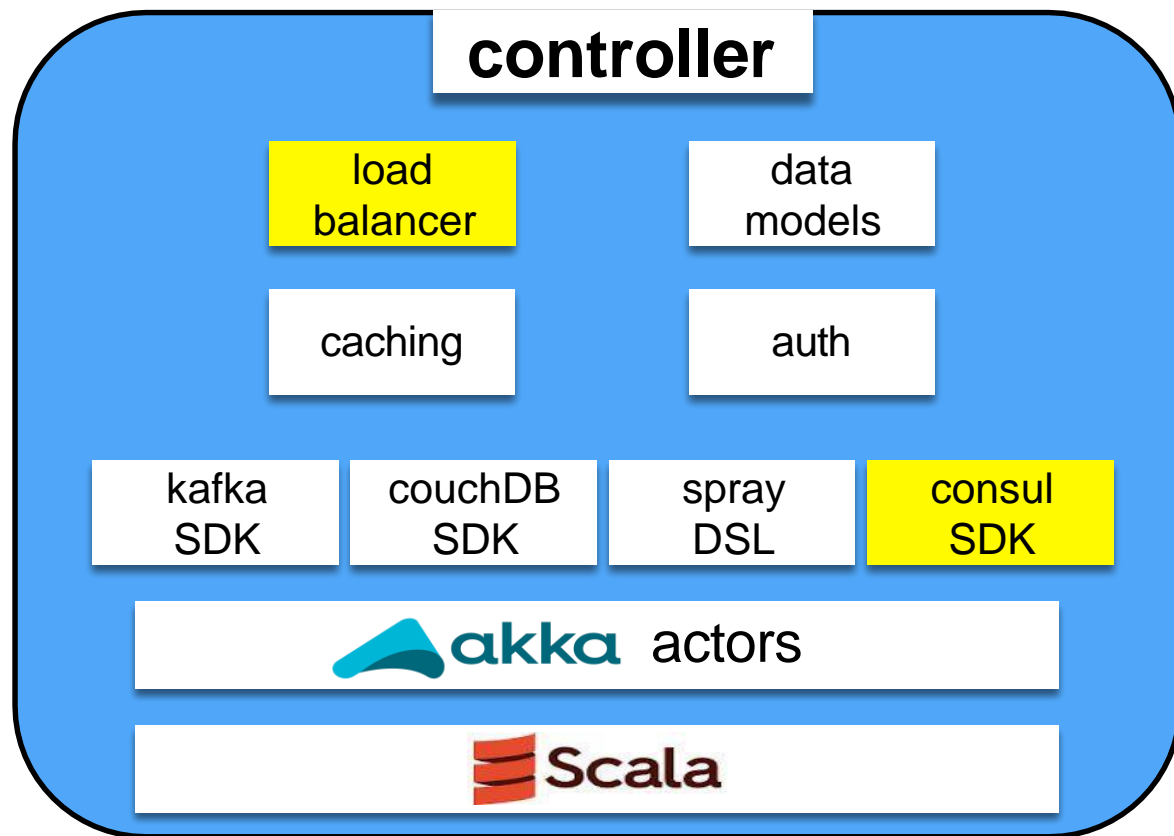


中国科学技术大学  
University of Science and Technology of China

- check resource limits
- actions stored as documents in CouchDB
  - binaries as objects (attachments)



Load balancer: find a slave to  
execute Slave health, load in



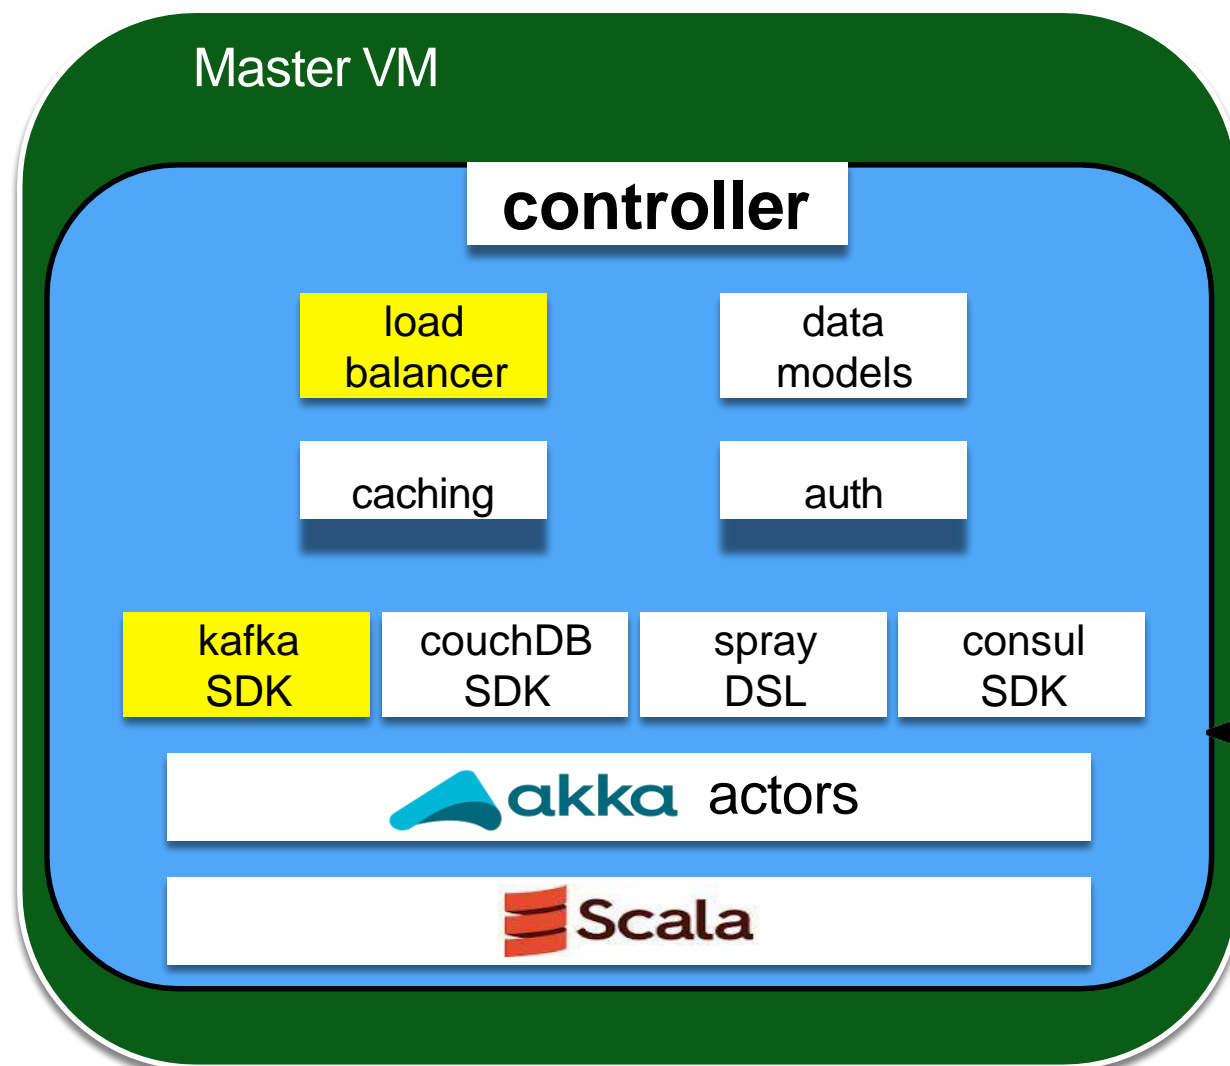
*Why* CONSUL ?

- Sequentially consistent KV store
- Replication, Fault Tolerance
- Health Check / Monitoring utilities

# Apache OpenWhisk: Step 6. Get in line!




Post request to execute to queue in



Why  kafka ?

- High throughput fault-tolerant queues
- *Point-to-point* messages via topics
  - explicit load balancing





# Apache OpenWhisk: Step7. Get to Work!



中国科学技术大学  
University of Science and Technology of China

Slave VM

**invoker**



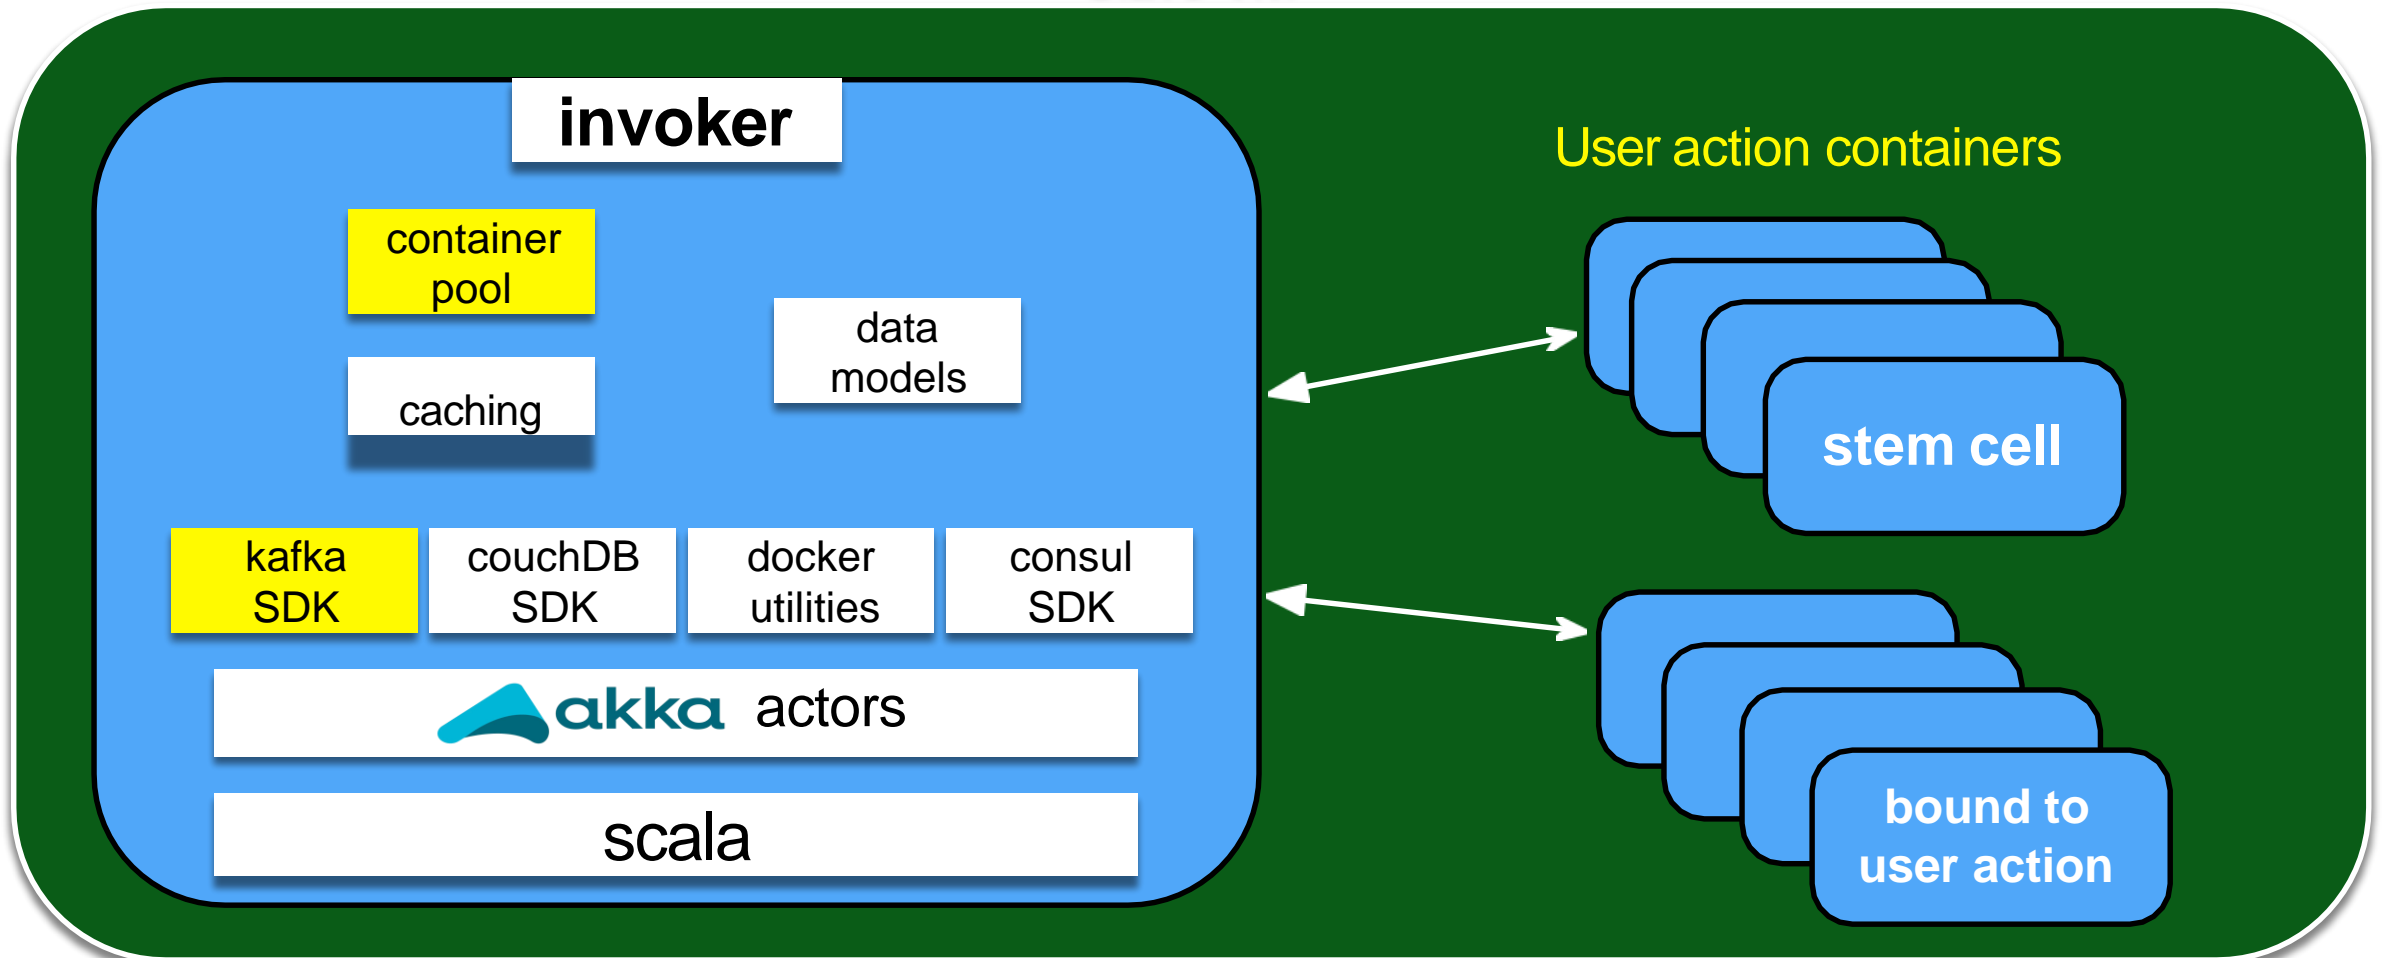
# Apache OpenWhisk: Step 7. Get to work!

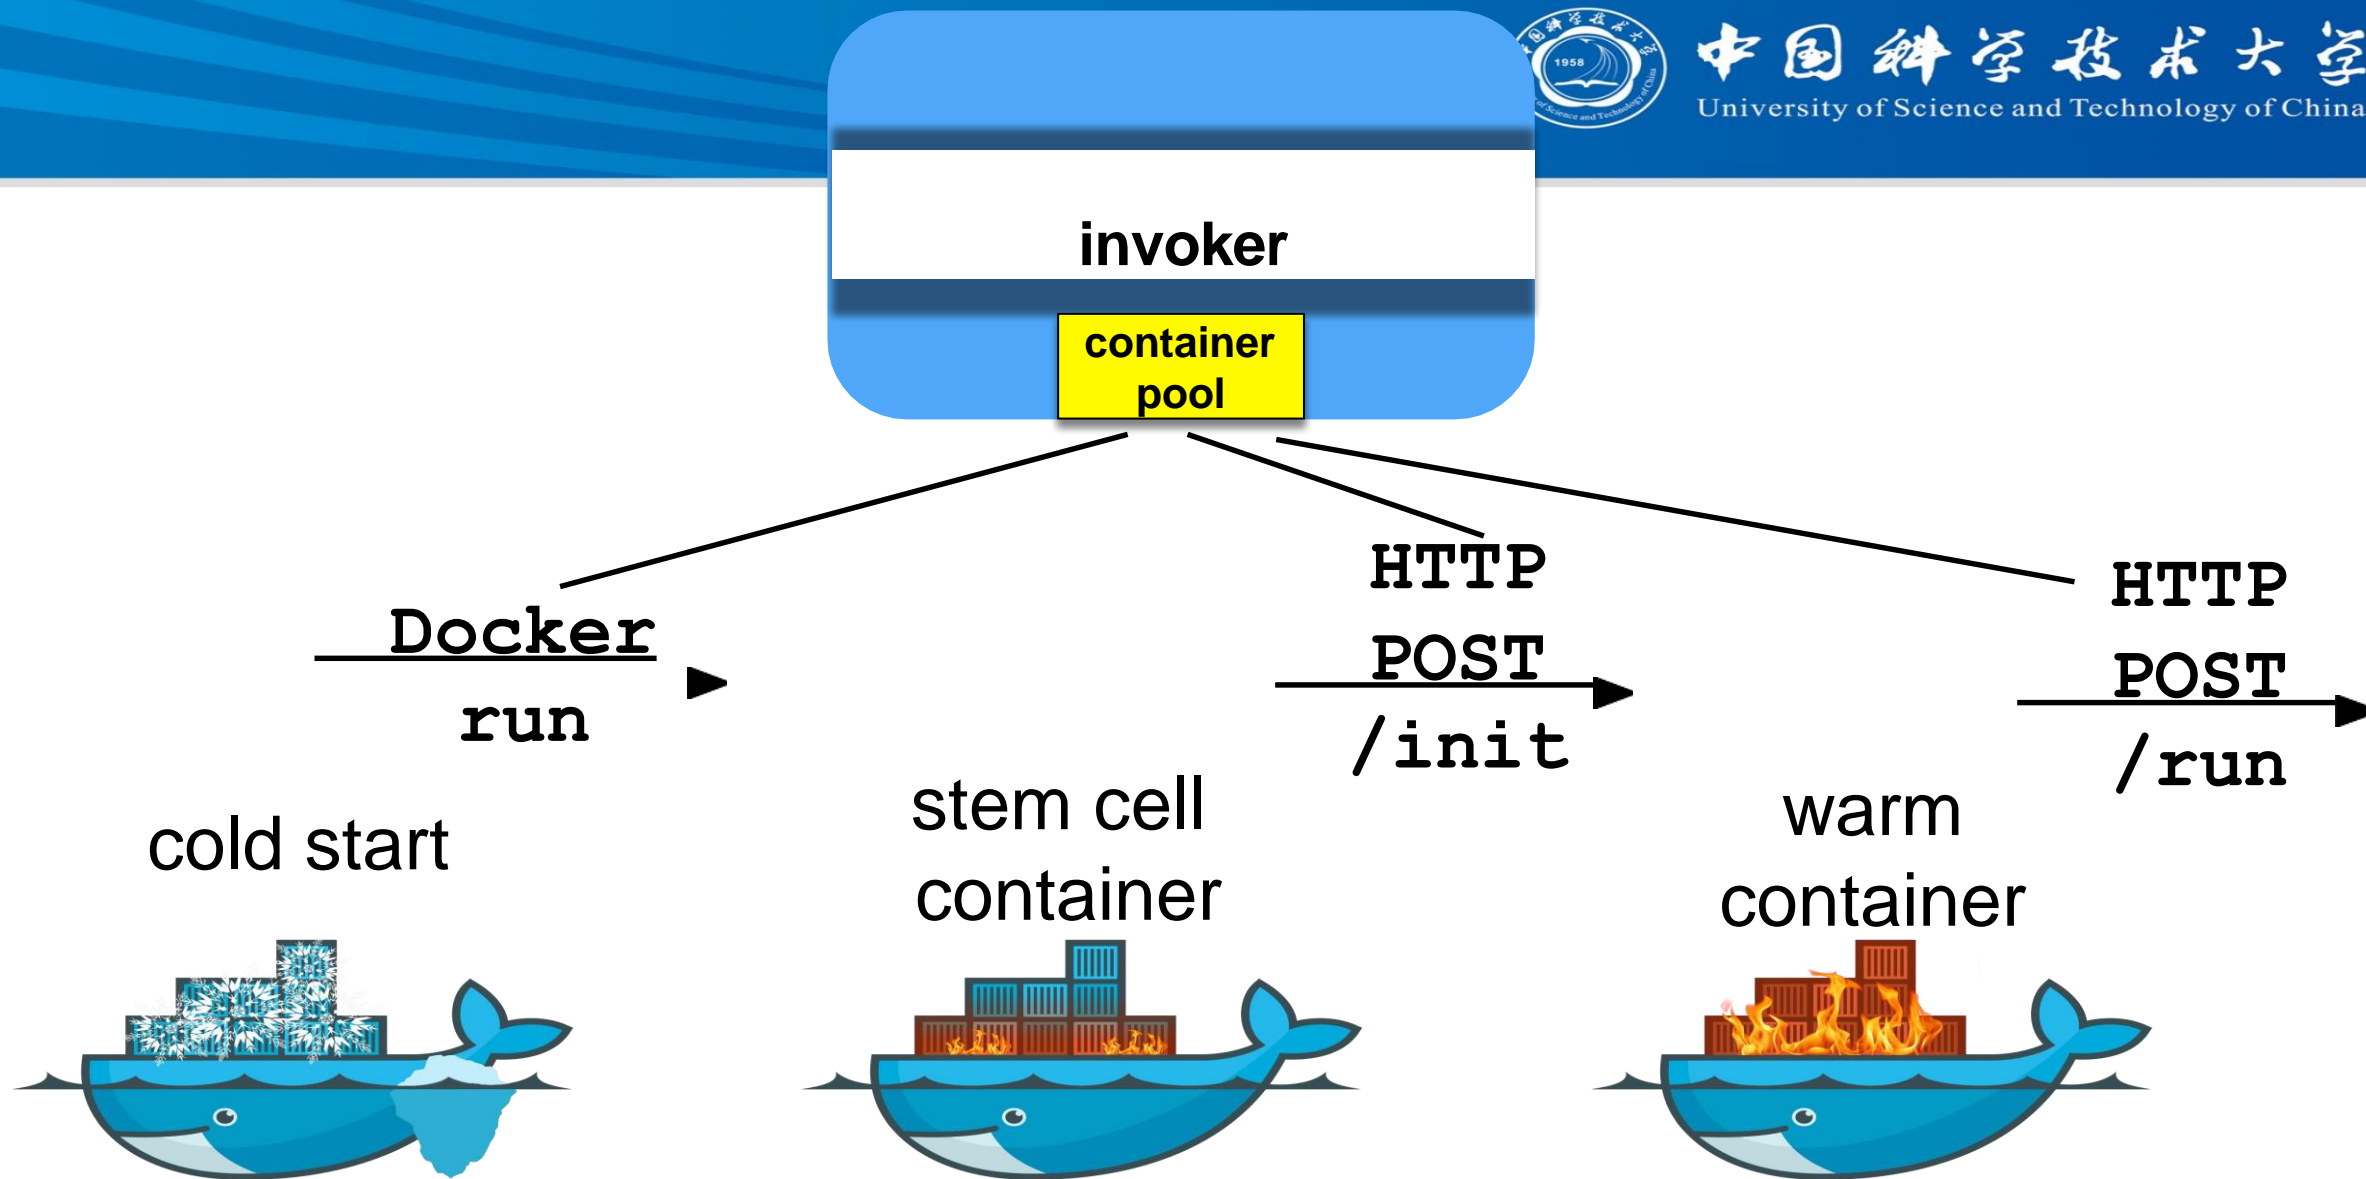


中国科学技术大学  
University of Science and Technology of China

- each user action gets its own container (isolation)
- containers may be reused
- container pool allocates and garbage collects containers

Slave VM

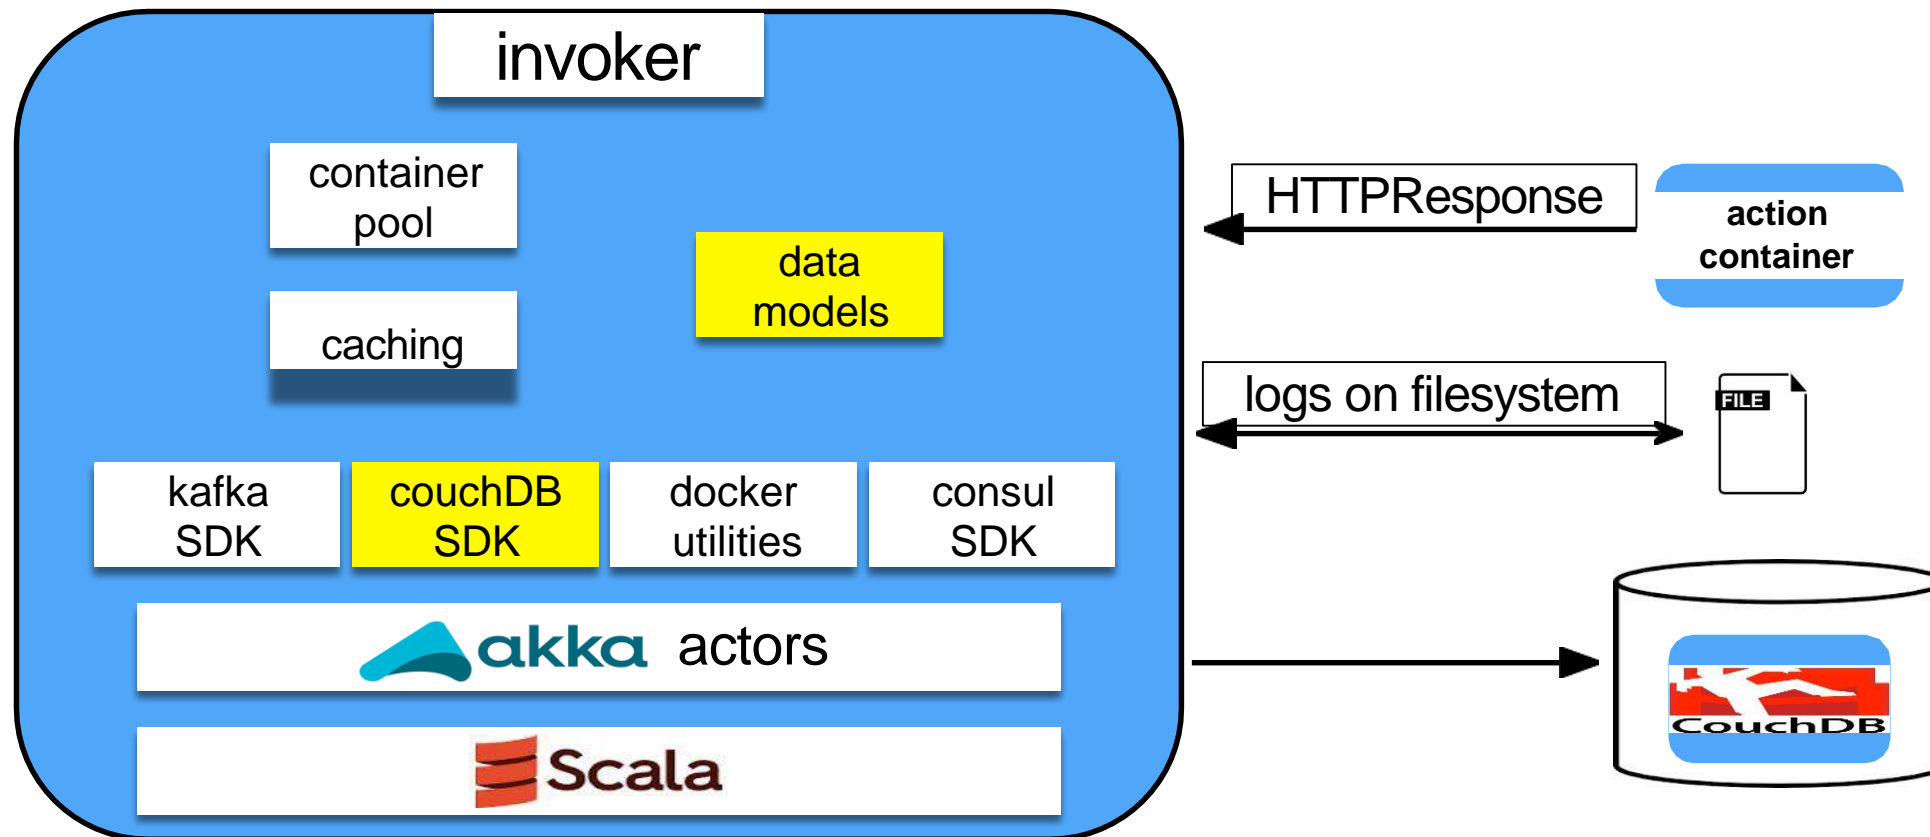




# Apache OpenWhisk: Step 8. Store the results.



中国科学技术大学  
University of Science and Technology of China





- Cold start problem
  - Keep invokers ready (“stem cell”) or running (“warm”) after invocation
  - Tradeoff with latency and resource reservation
- Auto scale
  - Add to and remove from the invoker pool
  - Hibernate when idle
- Fine-grained billing
  - Overhead of metering
  - Choice of which resources to bill (CPU, memory, network, ...)
  - Understandable billing policy (simple vs detailed)?




- Reactive programming
- Event-based applications
- Stream processing systems
- Dataflow programming
- Workflows and business processes
- Service composition
- Service oriented architectures
- many more ...



# Serverless Apps Lifecycle



中国科学技术大学  
University of Science and Technology of China





## Distributed Architecture

- Design stateless and ASync solutions to enable scaling.
- Connect with other Azure Services via triggers and bindings.
- Use Logic Apps to orchestrate workflows
- Use managed connectors to abstract calls to cloud and on-premises services.

## Cloud DevOps

- Design for automation. Use ARM templates.
- Design DevOps for the cloud: safe deployment with test/development and production environment separation and test on the target platform.
- Monitor the running apps with App Insights and tune for best experience.

# Triggers and Bindings



中国科学技术大学  
University of Science and Technology of China

## Triggers and Bindings

Type	Service	Trigger	Input	Output
Schedule	Azure Functions	✓		
HTTP (REST or webhook)	Azure Functions	✓		✓*
Blob Storage	Azure Storage	✓	✓	✓
Events	Azure Event Hubs	✓		✓
Queues	Azure Storage	✓		✓
Queues and topics	Azure Service Bus	✓		✓
Tables	Azure Storage		✓	✓
Tables	Azure Mobile Apps		✓	✓
No-SQL DB	Azure DocumentDB		✓	✓
Push Notifications	Azure Notification Hubs			✓
Twilio SMS Text	Twilio			✓

# Use Bindings in Your Code



function.json

```
public static class OrderHandler
{
    [FunctionName("OrderWebhook")]
    public static async Task<HttpResponseMessage> Run(
        [HttpTrigger] HttpRequestMessage req,
        [Queue("aievents1", Connection = "AiStorageConnection")]
            IAsyncCollector<String> eventOutput,
            TraceWriter log)
    {
        log.Info($"Webhook was triggered!");

        string jsonContent = await req.Content.ReadAsStringAsync();
        dynamic data = JsonConvert.DeserializeObject(jsonContent);

        await eventOutput.AddAsync(
            JsonConvert.SerializeObject(GetLogData(data)));

        int orderId = PlaceOrder(data);

        return req.CreateResponse(HttpStatusCode.OK,
            new {orderNumber = orderId });
    }
    . . .
}
```

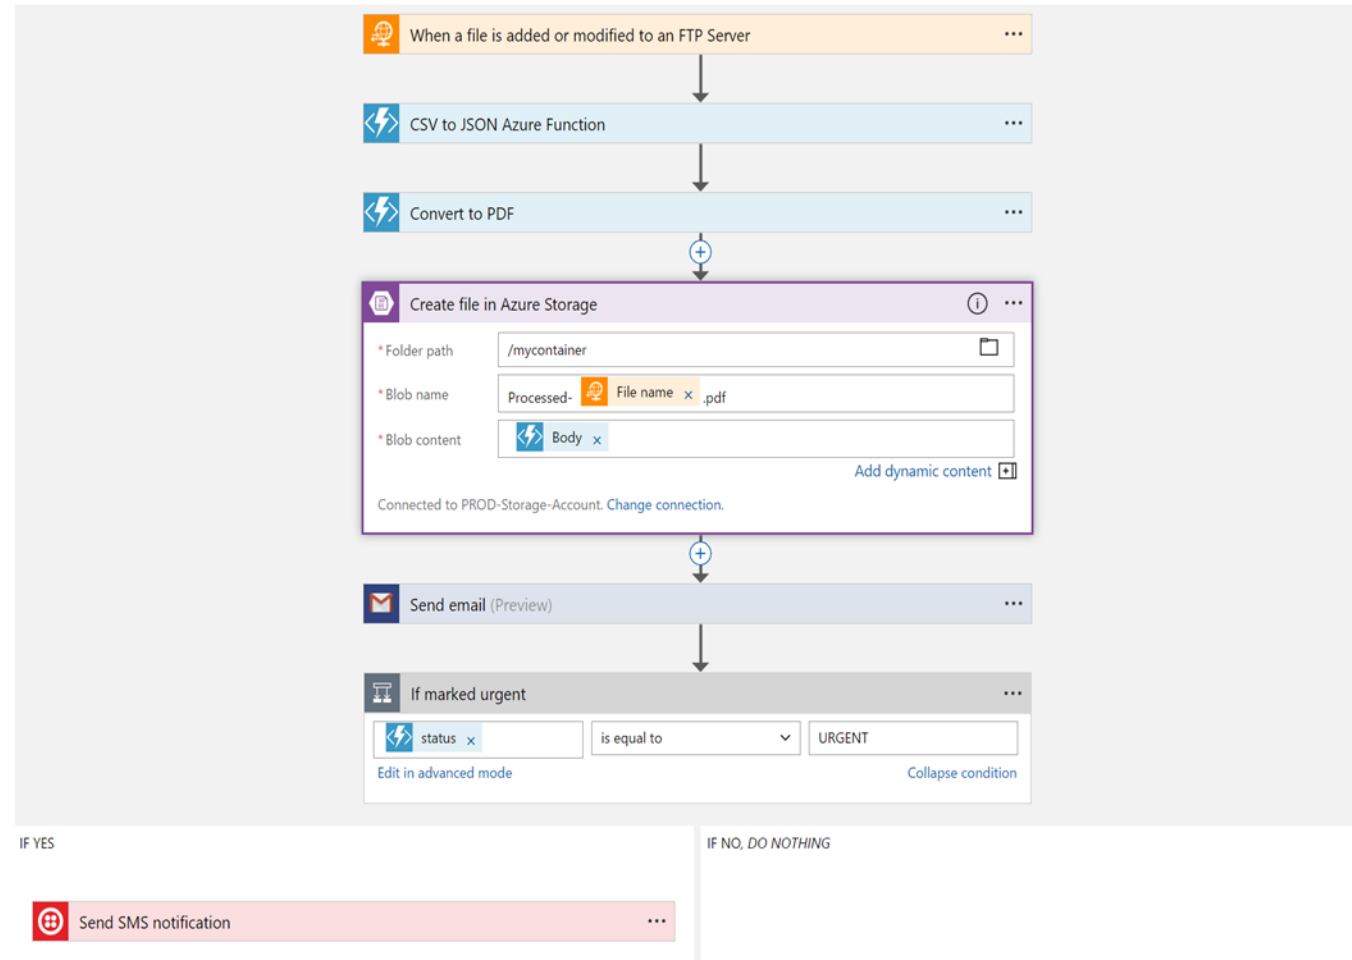
```
"bindings": [
  {
    "type": "httpTrigger",
    "direction": "in",
    "webHookType": "genericJson",
    "name": "req"
  },
  {
    "type": "http",
    "direction": "out",
    "name": "res"
  },
  {
    "type": "queue",
    "name": "eventOutput",
    "queueName": "aievents1",
    "connection": "AiStorageConnection",
    "direction": "out"
  }
]
```

# Logic Apps Workflow Designer



中国科学技术大学  
University of Science and Technology of China

- Workflow in the cloud
- Powerful control flow
- Connect functions and APIs
- Declarative definition to persist in source control and drive deployments



# Logic Apps

## SaaS

- appFigures
- Asana
- Azure API Management
- Azure App Services
- Azure Automation
- Azure Cognitive Face API
- Azure Cognitive LUIS
- Azure Cognitive Text Analytics
- Azure Cognitive Vision
- Azure Data Lake Store
- Azure Document DB
- Azure Event Hub
- Azure Functions
- Azure Machine Learning
- Azure Resource Manager
- Azure Service Bus
- Azure SQL
- Azure Storage Blob
- Azure Storage Queues
- Basecamp
- Bing Search
- BitBucket
- Bitly
- Blogger
- Box
- Buffer
- Campfire
- Chatter
- Common Data Service
- Disqus
- DocuSign
- Dropbox
- Dynamics AX Online
- Dynamics CRM Online
- Dynamics CRM Service Bus
- Dynamics Financials
- Dynamics Operations
- Easy Redmine
- Eventbrite
- Facebook
- FreshBooks
- Freshdesk
- GitHub
- Gmail
- Google Calendar
- Google Contacts
- Google Drive
- Google Sheets
- Google Tasks
- GoTo Meeting
- GoTo Training
- GoTo Webinar
- Harvest
- HelloSign
- Infusionsoft
- JIRA
- Insightly
- Instagram
- Instapaper
- MailChimp
- Mandrill
- Medium
- Microsoft Project Online
- Microsoft Translator
- MSN Weather
- Muhimbi PDF
- Office 365
- Office 365 Users
- Office 365 Video
- OneDrive
- OneDrive for Business
- OneNote
- Outlook.com
- Outlook Tasks
- PagerDuty
- Pinterest
- Pipedrive
- Pivotal Tracker
- Power BI
- Project Online
- Redmine
- Salesforce
- Salesforce Chatter
- SendGrid
- SharePoint Online
- Slack
- SmartSheet
- SparkPost
- Stripe
- Survey Monkey
- Todoist
- Toodledo
- Trello
- Twilio
- Twitter
- Typeform
- UserVoice
- VS Team Services
- Webmerge
- Wordpress
- Wunderlist
- Yammer
- YouTube
- Zendesk
- HTTP Webhook
- FTP, SFTP
- SMTP
- RSS
- Compose, Query, Parse JSON
- Wait
- Terminate
- Workflow

## Cloud APIs and platform

- Supports over 125 built-in connectors
- Scales to meet your needs
- Enables rapid development
- Extends with custom APIs and Functions

## API connections

- Authenticate once and reuse

## XML and EDI

- XML Validation
- Transform XML (+Mapper)
- Flat File Encode
- Flat File Decode
- X12
- EDIFACT
- AS2
- Integration Account Artifact Lookup

## Hybrid

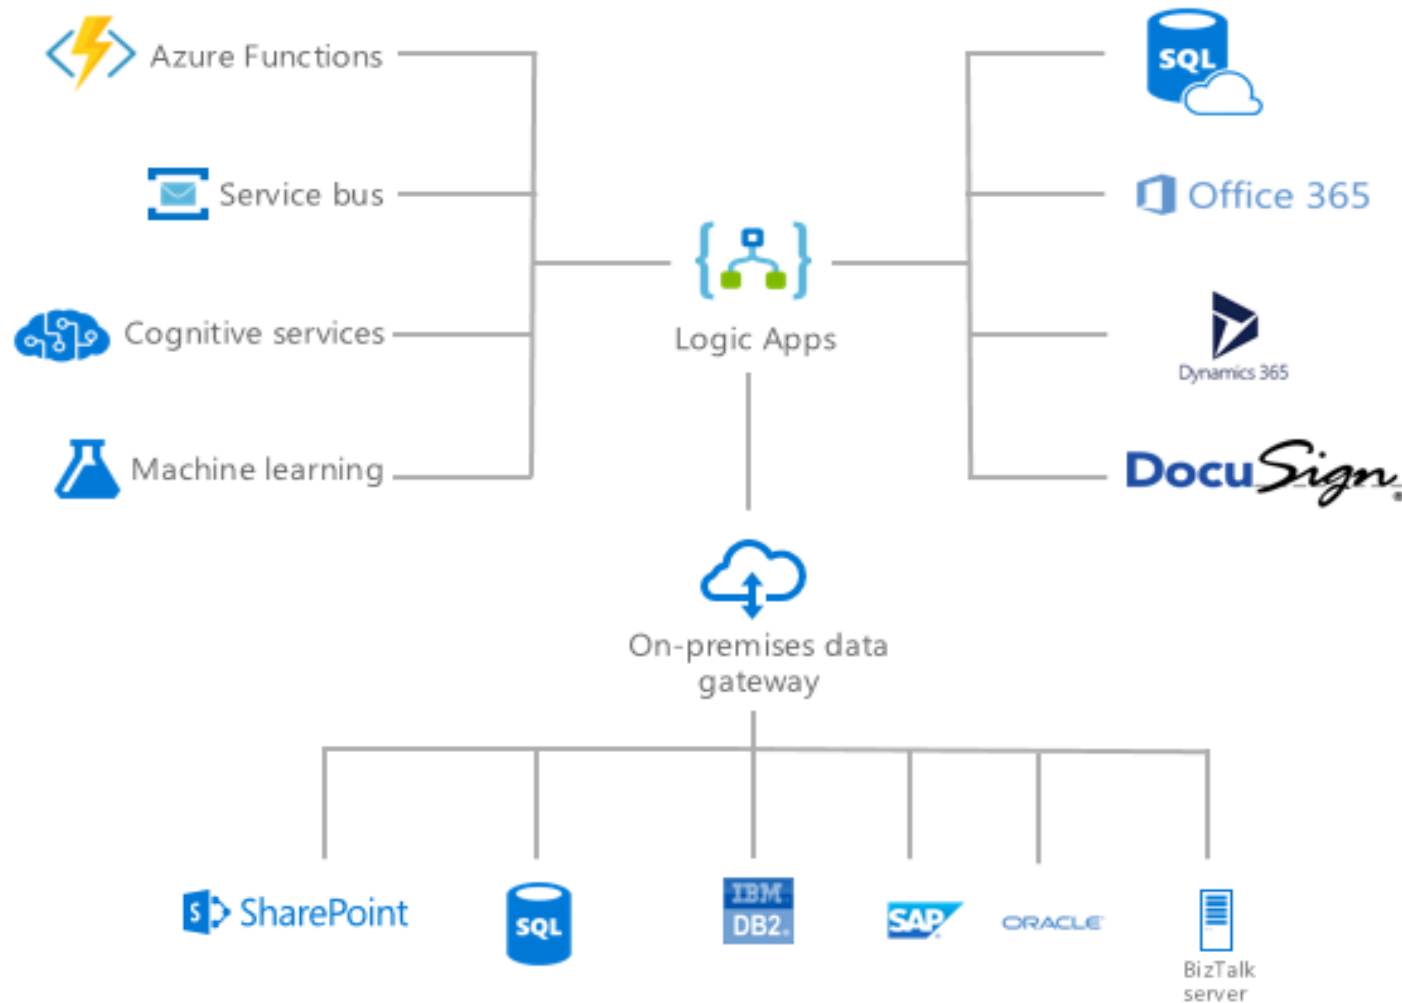
- BizTalk Server
- File System
- IBM DB2
- Informix
- Oracle DB
- SharePoint Server
- SQL Server
- SAP
- Websphere MQ



# Logic Apps Connect Everything



中国科学技术大学  
University of Science and Technology of China






# Develop



中国科学技术大学  
University of Science and Technology of China





## Azure Functions Core Tools

- Provides the entire Functions runtime

- Trigger off of Azure events and debug locally

## JavaScript

- Use Visual Studio Code or any Node debugger

## C#

- Use Visual Studio 2015 or 2017

- Use class libraries with attributes in Visual Studio 2017

# C# and Visual Studio



中国科学技术大学  
University of Science and Technology of China

- Based on class libraries
- Get the full power of IntelliSense, unit testing, and local debugging
- Use Web Jobs attributes to define triggers and bindings

The screenshot shows the Visual Studio IDE with a C# file named `OrderHandler.cs` open. The code defines a `public static class OrderHandler` with a `Run` method. The `Run` method is an `async Task` that takes an `HttpRequestMessage req` and an `IAsyncCollector<String> eventOutput` as parameters. It logs a message, reads the request content as a string, deserializes it into a `dynamic` object, adds the serialized object to the event output, places an order, and returns a response with the order ID.

```
15 public static class OrderHandler
16 {
17     [FunctionName("OrderWebhook")]
18     public static async Task<HttpResponseMessage> Run(
19         [HttpTrigger] HttpRequestMessage req,
20         [Queue("aievents1", Connection = "AiStorageConnection")] IAsyncCollector<String> eventOutput,
21         TraceWriter log)
22     {
23         log.Info($"Webhook was triggered!");
24
25         string jsonContent = await req.Content.ReadAsStringAsync();
26         dynamic data = JsonConvert.DeserializeObject(jsonContent);
27
28         await eventOutput.AddAsync(
29             JsonConvert.SerializeObject(GetLogData(data)));
30
31         int orderId = PlaceOrder(data); ≤ 727ms elapsed
32
33         return req.CreateResponse(HttpStatusCode.OK, new {orderNumber = orderId });
34     }
35 }
```

The Watch window at the bottom shows the following data:


Name	Value	Type
((dynamic)data).email="builddebug@test.com"	"builddebug@test.com"	dynamic
data	{ "time": "2017-05-09T03:26:09.407Z", "email": "builddebug@test.com", "order": [ { "type": "Margherita", "size": object { }	dynamic

Learn more at <https://aka.ms/vs2017functiontools>

# Deploy



中国科学技术大学  
University of Science and Technology of China



# Deployment Options



中国科学技术大学  
University of Science and Technology of China

## Resource deployment

- Azure Resource Manager (i.e. ARM)

## Content deployment

- Visual Studio
- Azure CLI (Logic App)
- Azure Functions Core Tools (Function App)
- CI/CD

<https://www.visualstudio.com/en-us/docs/build/get-started/aspnet-4-ci-cd-azure-automatic>


## Safe deployment practices

- Use Azure Functions deployment slots for environment separation and swap deployments

# Monitor




中国科学技术大学  
University of Science and Technology of China



# Key Scenarios for Monitoring




中国科学技术大学  
University of Science and Technology of China







- Extensible Application Performance Management (APM)
- Rich data: Metrics, Traces, Exception tracking, Dependencies, Page Views, User data, custom events
- Easy to use graph/alerts, powerful analytics portal, integration with PowerBI and other analytics services





## Developer experience

- Same consistent Programming Model
- Same Azure Functions portal
- Publish directly from Visual Studio tooling
- Leverage triggers: timer trigger and new SQL Service Broker trigger

## Administrator features

- Take advantage of Azure Functions on premises
  - Workers can run in spare compute – i.e. desktops left on overnight within orgs
- Only provision two types of roles
  - Management Role – Hosts Portal, Publishing Endpoint and
  - Worker Role – Runs Function in Windows Server Containers

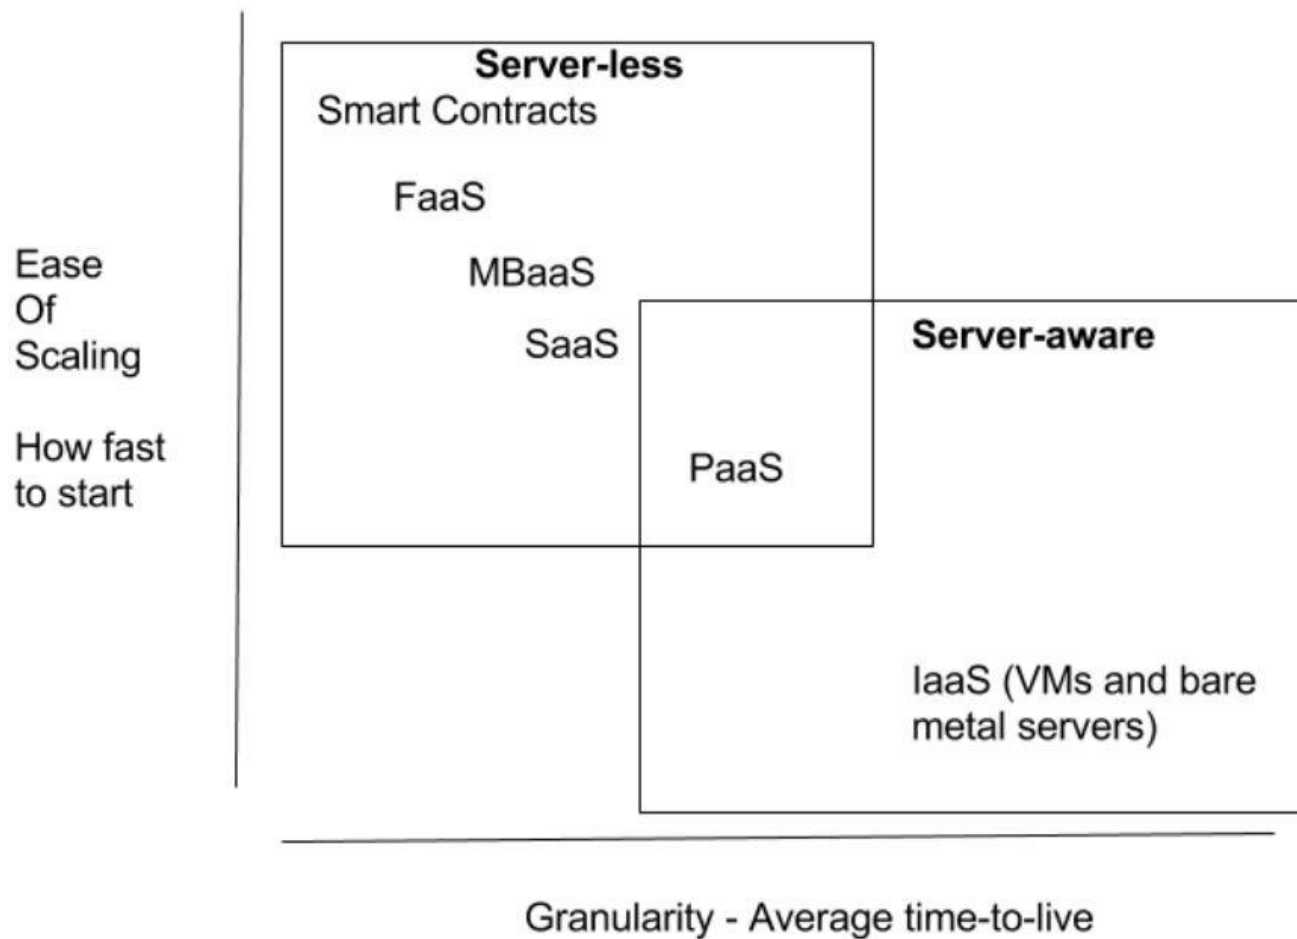




- Cost - pay-as-you-go is enough?
- Server-less - can servers be really hidden?
- Problem of state: stateless, state in other place, or state-ful supported in FaaS?
- Security - no servers!
- Legacy systems and serverless?
  - Hybrid model?



# Cloud computing: server-less vs server-aware?





- Tools
- Deployment
- Monitoring and debugging
  - Short-lived functions, scaling to large invocations,
  - Looking for problems is like finding needles in ever growing haystack?
- Serverless IDEs?
- Decompose micro-service into FaaS?
  - Code granularity is function?
- Managing state inside and outside FaaS
- Concurrency, recovery semantics, transactions?

# Open Problems - how FaaS fits into cloud?



中国科学技术大学  
University of Science and Technology of China

- Just another \*aaS?
- Can different cloud computing service models be mixed?
- Can there be more choices for how much memory and CPU can be used by serverless functions?
- Does serverless need to have IaaS-like based pricing?
- What about spot and dynamic pricing with dynamically changing granularity?





- Granularity of serverless is much smaller than traditional server based tool
- Debugging is much different if instead of having one artifact (a micro-service or traditional monolithic app) developers need to deal with a myriad of smaller pieces of code ...
  - That haystack can grow really big really fast ...

# Open Problems: can “legacy” code be made to run serverless



中国科学技术大学  
University of Science and Technology of China

- Today the amount of existing ( “legacy” ) code that must continue running is much larger than the new code created specifically to run in serverless environments
- The economical value of existing code represents a huge investment of countless hours of developers coding and fixing software
- Therefore, one of the most important problems may be to what degree existing legacy code can be automatically or semi-automatically decomposed into smaller-granularity pieces to take advantage of these new economics?



- Is serverless fundamentally stateless?
- Current serverless platforms are stateless will there be stateful serverless services in future?
- Will there be simple ways to deal with state?
- Can there be serverless services that have stateful support built-in
  - And with different degrees of quality-of-service?



- Combine low granularity basic building blocks of serverless (functions, actions, triggers, packages, ...) into bigger solutions?
- How to decompose apps into functions so that they use resources optimally?
- Are there lessons learned that can be applied from OOP design patterns, Enterprise Integration Patterns, etc.?



- IF functions is running outside of data-center is it serverless?
  - Cost, scalability, ...
- Internet of Things (IoT) will have many small devices each capable of running small amount of code - like functions in serverless?
- New domains, new concerns?
  - For example for IoT energy usage may be more important than speed?
- Are Blockchain smart contracts server-less?
  - For example when Ethereum users are running smart contracts they get paid for the “gas” consumed by the code, similar to fuel cost for an automobile but applied to computing (no need for data-center!)